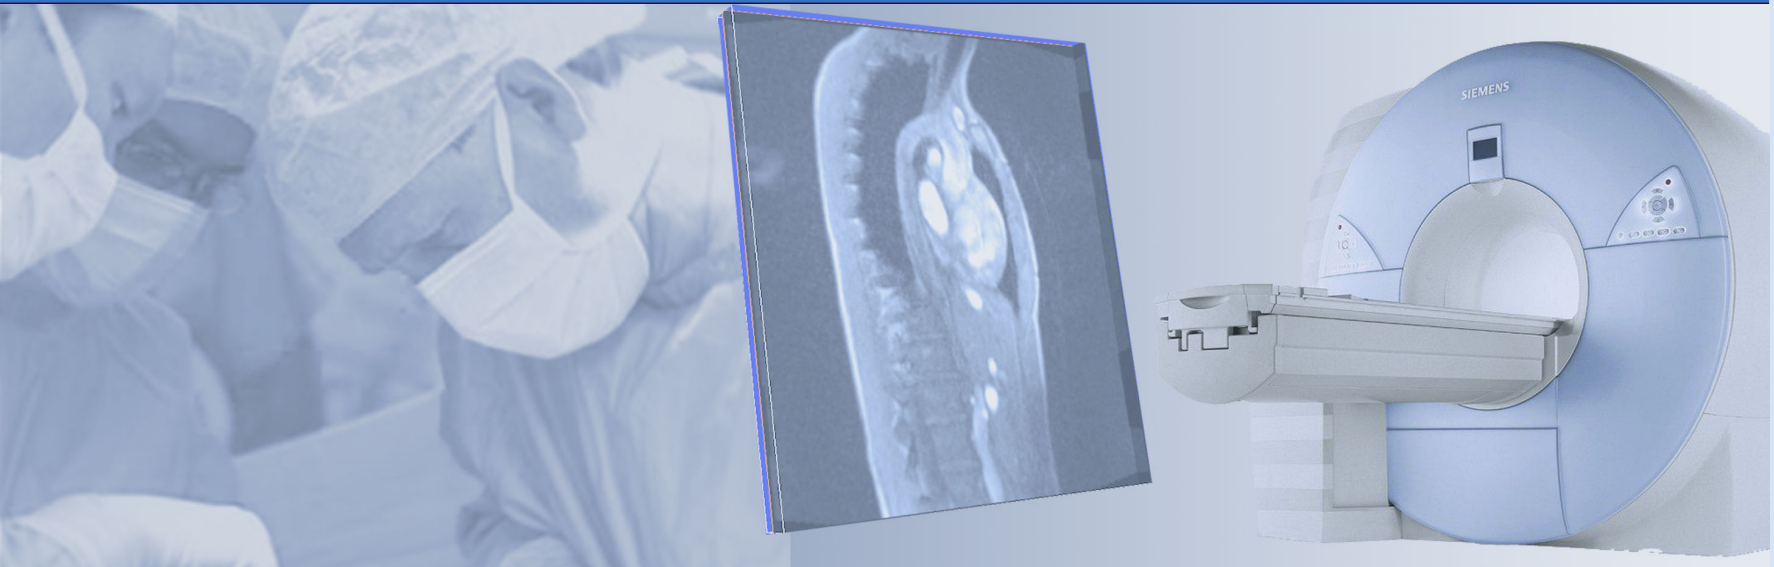


Tutorial Computer- and robot-assisted Surgery



NATIONALES CENTRUM
FÜR TUMORERKRANKUNGEN
PARTNERSTANDORT DRESDEN
UNIVERSITÄTS KREBSCENTRUM UCC

getragen von:

Deutsches Krebsforschungszentrum
Universitätsklinikum Carl Gustav Carus Dresden
Medizinische Fakultät Carl Gustav Carus, TU Dresden
Helmholtz-Zentrum Dresden-Rossendorf

Sebastian Bodenstedt
Translational Surgical Oncology

Questions lecture

Evaluation Metrics

Evaluation metrics

- Accuracy: how many samples were correctly identified?
 - $\frac{\text{\#correctly identified samples}}{\text{\#samples}}$
 - Works well for balanced data sets, i.e. all classes occurs similarly often
 - Problems with minor classes, e.g. diseases detection:
if 99% are negative examples, predicting just the negative classes leads to a high accuracy, but misses critical cases
- Confusion matrix

n=165	Predicted: NO	Predicted: YES
	Actual: NO	Actual: YES
	50	10
	5	100

Evaluation metrics (binary problems)

- Terms
 - True positives (TP): number of correctly identified positive examples
 - False positives (FP): number of falsely identified positive examples
 - True negatives (TN): number of correctly identified negative examples
 - False negatives (FN): number of falsely identified negative examples

		Predicted	
		Negative	Positive
Actual	Negative	True Negative	False Positive
	Positive	False Negative	True Positive

Evaluation metrics (binary problems)

- Precision
 - What portion of the as positive identified samples was correct?

$$\text{Precision} = \frac{\text{True positives}}{\text{True positives} + \text{False positives}}$$

		Predicted	
		Negative	Positive
Actual	Negative	True Negative	False Positive
	Positive	False Negative	True Positive

- Recall
 - What portion of the positive samples was found?

$$\text{Recall} = \frac{\text{True positives}}{\text{True positives} + \text{False negatives}}$$

		Predicted	
		Negative	Positive
Actual	Negative	True Negative	False Positive
	Positive	False Negative	True Positive

Evaluation metrics (binary problems)

- F1-score
 - Combination of precision and recall
 - How balanced are the two for your classifier?

$$F1\text{-score} = \frac{2 \times \text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

- Why use F1 instead of accuracy?
 - True negatives have no impact
- Binary metrics for multi-class problems
 - Generally binary metrics can be computed for each class
 - Positive samples vs negative samples for that class
 - The metrics for each class can then be aggregated, e.g. using mean or median

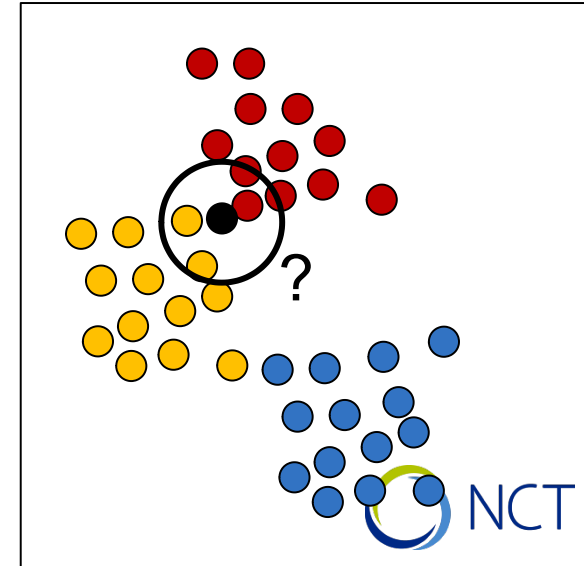
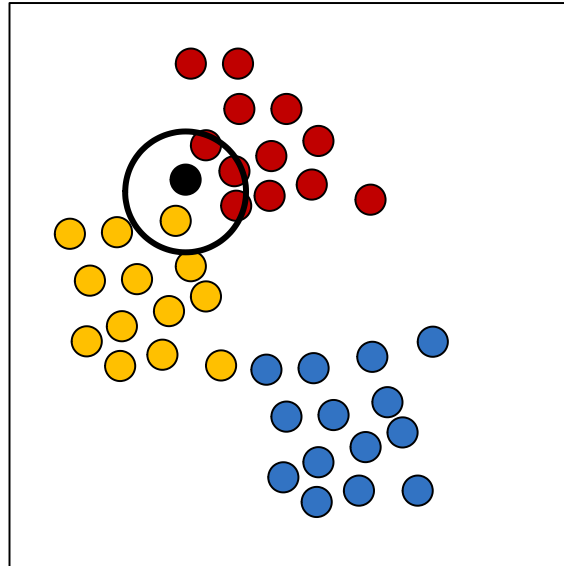
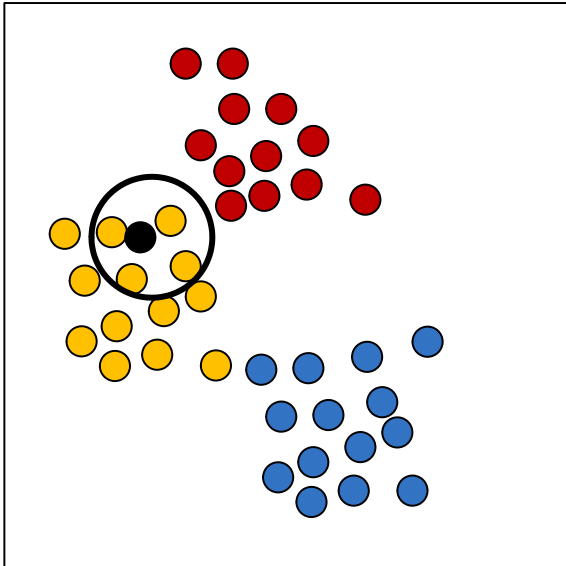
K-Nearest Neighbor

K-Nearest Neighbors

Idea

- Birds of a feather flock together
=> samples with similar values belong to the same label
- Instance-based learning, no generalization

$K = 4$



Principal Component Analysis (PCA)

Singular Value Decomposition (SVD)

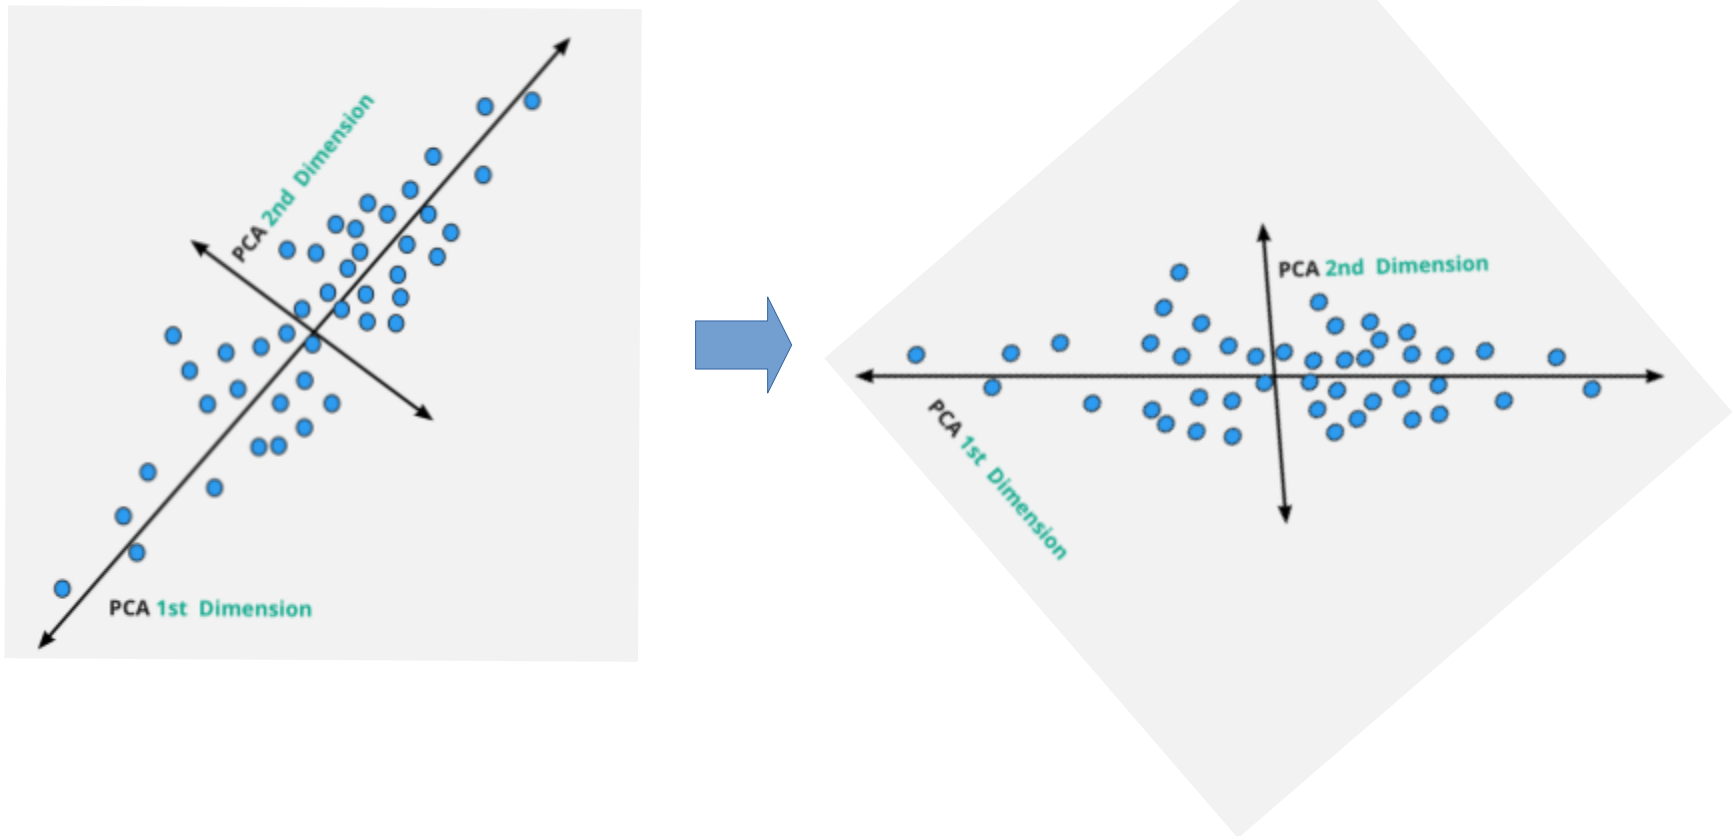
- Factorization of a $m \times n$ matrix into $A = U \Sigma V^T$
 - U : $m \times m$ orthonormal matrix. Contains eigenvectors of AA^T
 - V : $n \times n$ orthonormal matrix. Contains eigenvectors of $A^T A$
 - Σ : $m \times n$ diagonal matrix. Contains eigenvalues of $A^T A$

$$A = \begin{pmatrix} \vdots & \dots & \vdots \\ \mathbf{u}_1 & \dots & \mathbf{u}_n \\ \vdots & \dots & \vdots \end{pmatrix} \begin{pmatrix} \sigma_1 & & \\ & \ddots & \\ & & \sigma_n \end{pmatrix} \begin{pmatrix} \dots & \mathbf{v}_1^T & \dots \\ \vdots & \vdots & \vdots \\ \dots & \mathbf{v}_n^T & \dots \end{pmatrix}$$
$$A = \begin{pmatrix} \vdots & \dots & \vdots \\ \mathbf{u}_1 & \dots & \mathbf{u}_n \\ \vdots & \dots & \vdots \end{pmatrix} \begin{pmatrix} \sigma_1 & & \\ & \ddots & \\ & & \sigma_n \end{pmatrix} \begin{pmatrix} \vdots & \dots & \vdots \\ \mathbf{v}_1 & \dots & \mathbf{v}_n \\ \vdots & \dots & \vdots \end{pmatrix}^T$$

Reminder eigenvector and eigenvalue: $A\mathbf{v} = \lambda \mathbf{v}$

PCA

- Curse of dimensionality => number of required training samples increase (possibly exponentially) with the number of dimensions
- Idea: Re-orient data to maximize variance along axis and remove low-variance dimensions



PCA

- Steps

- 1) Normalize data

$$z = \frac{\text{value} - \text{mean}}{\text{standard deviation}}$$

- 2) Calculate covariance matrix (describes spread of the data)

$$\Sigma = \frac{1}{N} D \cdot D^T$$

- 3) Compute eigenvectors and eigenvalues of covariance matrix (via SVD)

- Eigenvectors are principal axis
- Eigenvalues indicate variance or “data spread” along axis

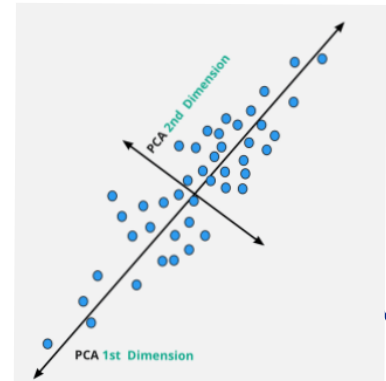
- 4) Build rotation matrix from principal axis (basis change)

- 5) Rotate data and remove dimensions with low variance

Retain the first t components that describe e.g. 98% of variation

$$V_T = \sum \lambda_i$$

$$\sum_{i=1}^t \lambda_i \geq f_v V_T$$



Introduction: Machine Learning with Scikit-learn

Scikit-learn

- Machine-learning library for Python
 - Includes many machine learning methods
 - Clustering
 - Random Forest
 - Support Vector Machine
 - ...
 - Uses numpy for most computations



Scikit-learn

- Linear classifier
 - Initialize model
`clf = SVC(kernel="linear")`
 - Fit model to data (x is data and y the corresponding labels)
`clf.fit(x, y)`
 - Predict label for data
`pred_train = clf.predict(x)`
`pred_test = clf.predict(x_test)`

Scikit-learn

- K-nearest neighbor
 - Initialize model (n_neighbors: number of neighbors to consider)
`clf = KNeighborsClassifier(n_neighbors=3)`
 - Fit model to data (x is data and y the corresponding labels)
`clf.fit(x, y)`
 - Predict label for data
`pred_train = clf.predict(x)`
`pred_test = clf.predict(x_test)`

Scikit-learn

- Decision tree
 - Initialize model (max_depth: how deep is the tree allowed to become)
`clf = DecisionTreeClassifier(random_state=0,
max_depth=5)`
 - Fit model to data (x is data and y the corresponding labels)
`clf.fit(x, y)`
 - Predict label for data
`pred_train = clf.predict(x)`
`pred_test = clf.predict(x_test)`

Scikit-learn

- Random forest
 - Initialize model (max_depth: how deep is the tree allowed to become, n_estimators: how many trees should the ensemble contain)
`clf = RandomForestClassifier(max_depth=5,
random_state=0, criterion="entropy", n_estimators=100)`
 - Fit model to data (x is data and y the corresponding labels)
`clf.fit(x, y)`
 - Predict label for data
`pred_train = clf.predict(x)`
`pred_test = clf.predict(x_test)`

Scikit-learn

- PCA
 - Initialize model (n_components: how many components should the PCA retain)
`pca = PCA(n_components=4)`
 - Fit model to data (x is data and y the corresponding labels)
`pca.fit(x)`
 - Predict label for data
`x = pca.transform(x)`
`x_test = pca.transform(x_test)`

**Are there
any
questions
?**