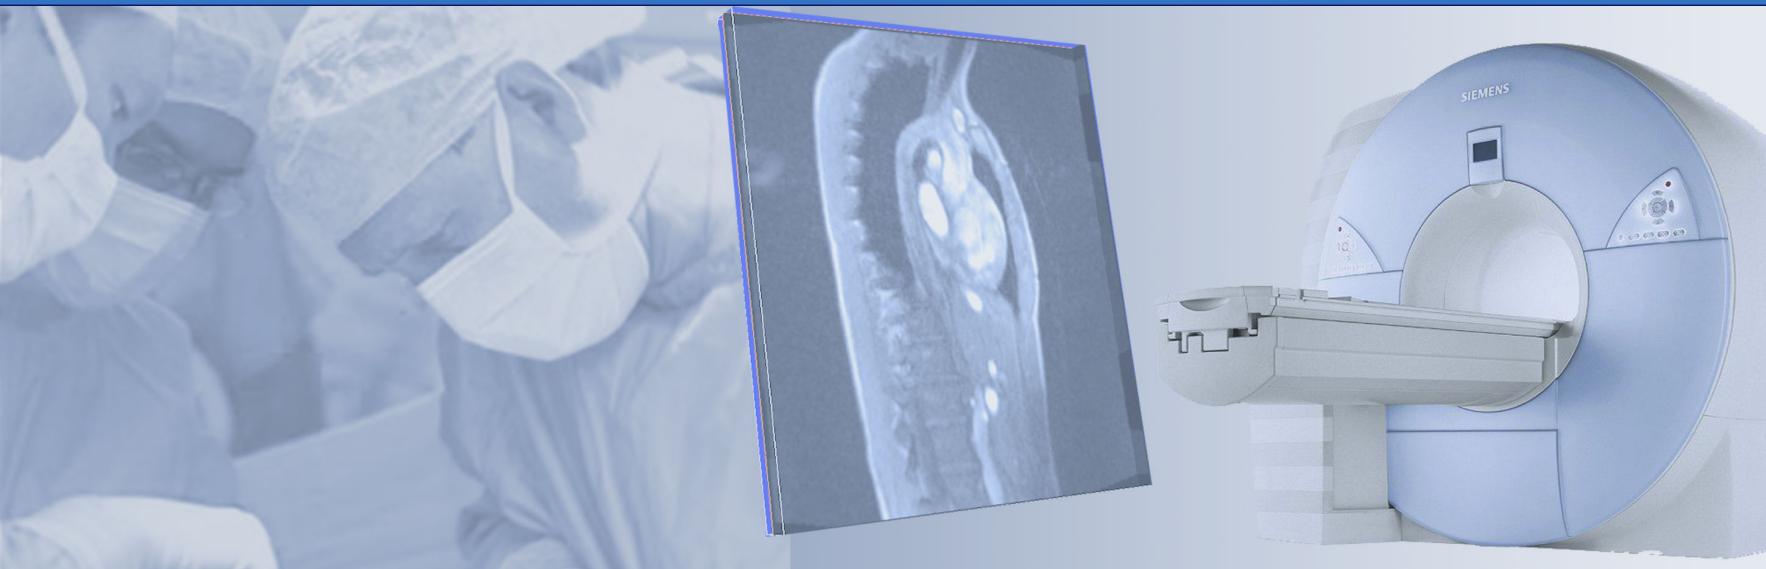


Computer- and robot-assisted Surgery



NATIONALES CENTRUM
FÜR TUMORERKRANKUNGEN
PARTNERSTANDORT DRESDEN
UNIVERSITÄTS KREBSZENTRUM UCC

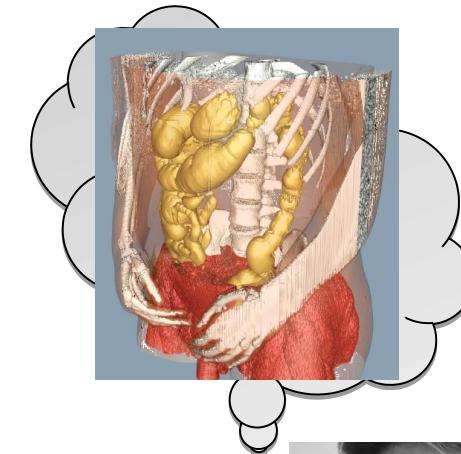
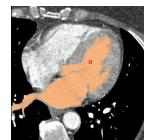
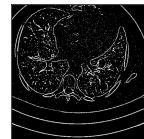
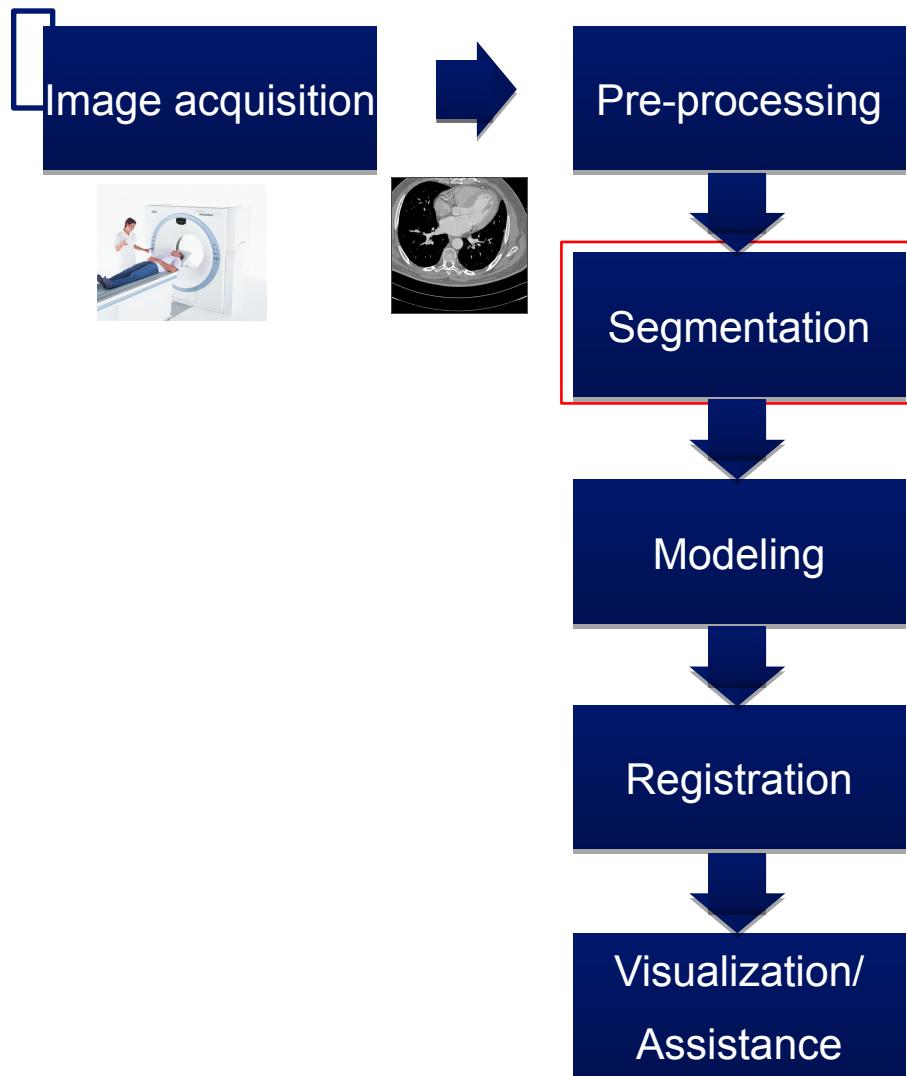
getragen von:
Deutsches Krebsforschungszentrum
Universitätsklinikum Carl Gustav Carus Dresden
Medizinische Fakultät Carl Gustav Carus, TU Dresden
Helmholtz-Zentrum Dresden-Rossendorf

Sebastian Bodenstedt
Lecture 10 – Segmentation III/ML Basics II

Oral Exam

If you want to participate in the oral exam, please contact
Ricarda Abdel-Bary (ricarda.abdel-bary@nct-dresden.de)

Process chain

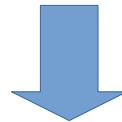


Contents

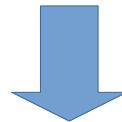
Neural Networks: How can images be classified and segmented?

- What are artificial neural networks?
 - How can I train a neural network?
- Convolutional Neural Networks
 - Image classification (what can be seen in an image)
 - Image segmentation (where can something be seen in an image)

Image classification: Which objects are in an image?



Classifier



Dog or cat?

Image classification: Which objects are in an image?



Image segmentation: Where in the image are objects?

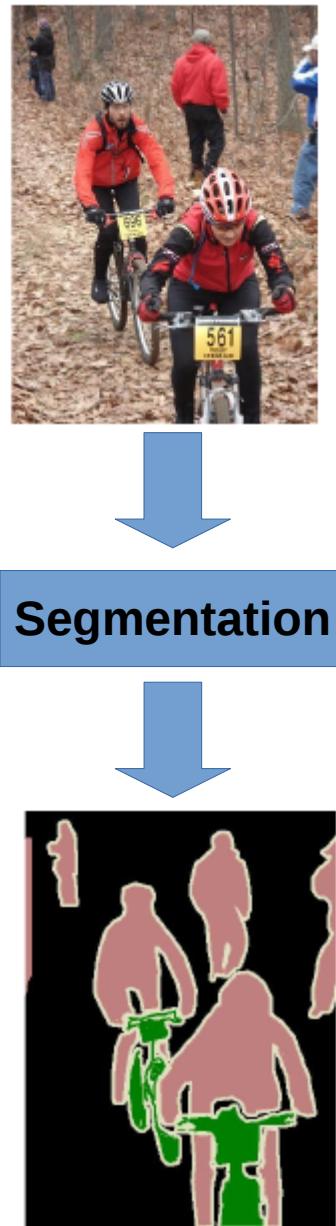
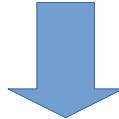
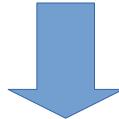


Image segmentation: Where in the image are objects?



Segmentation



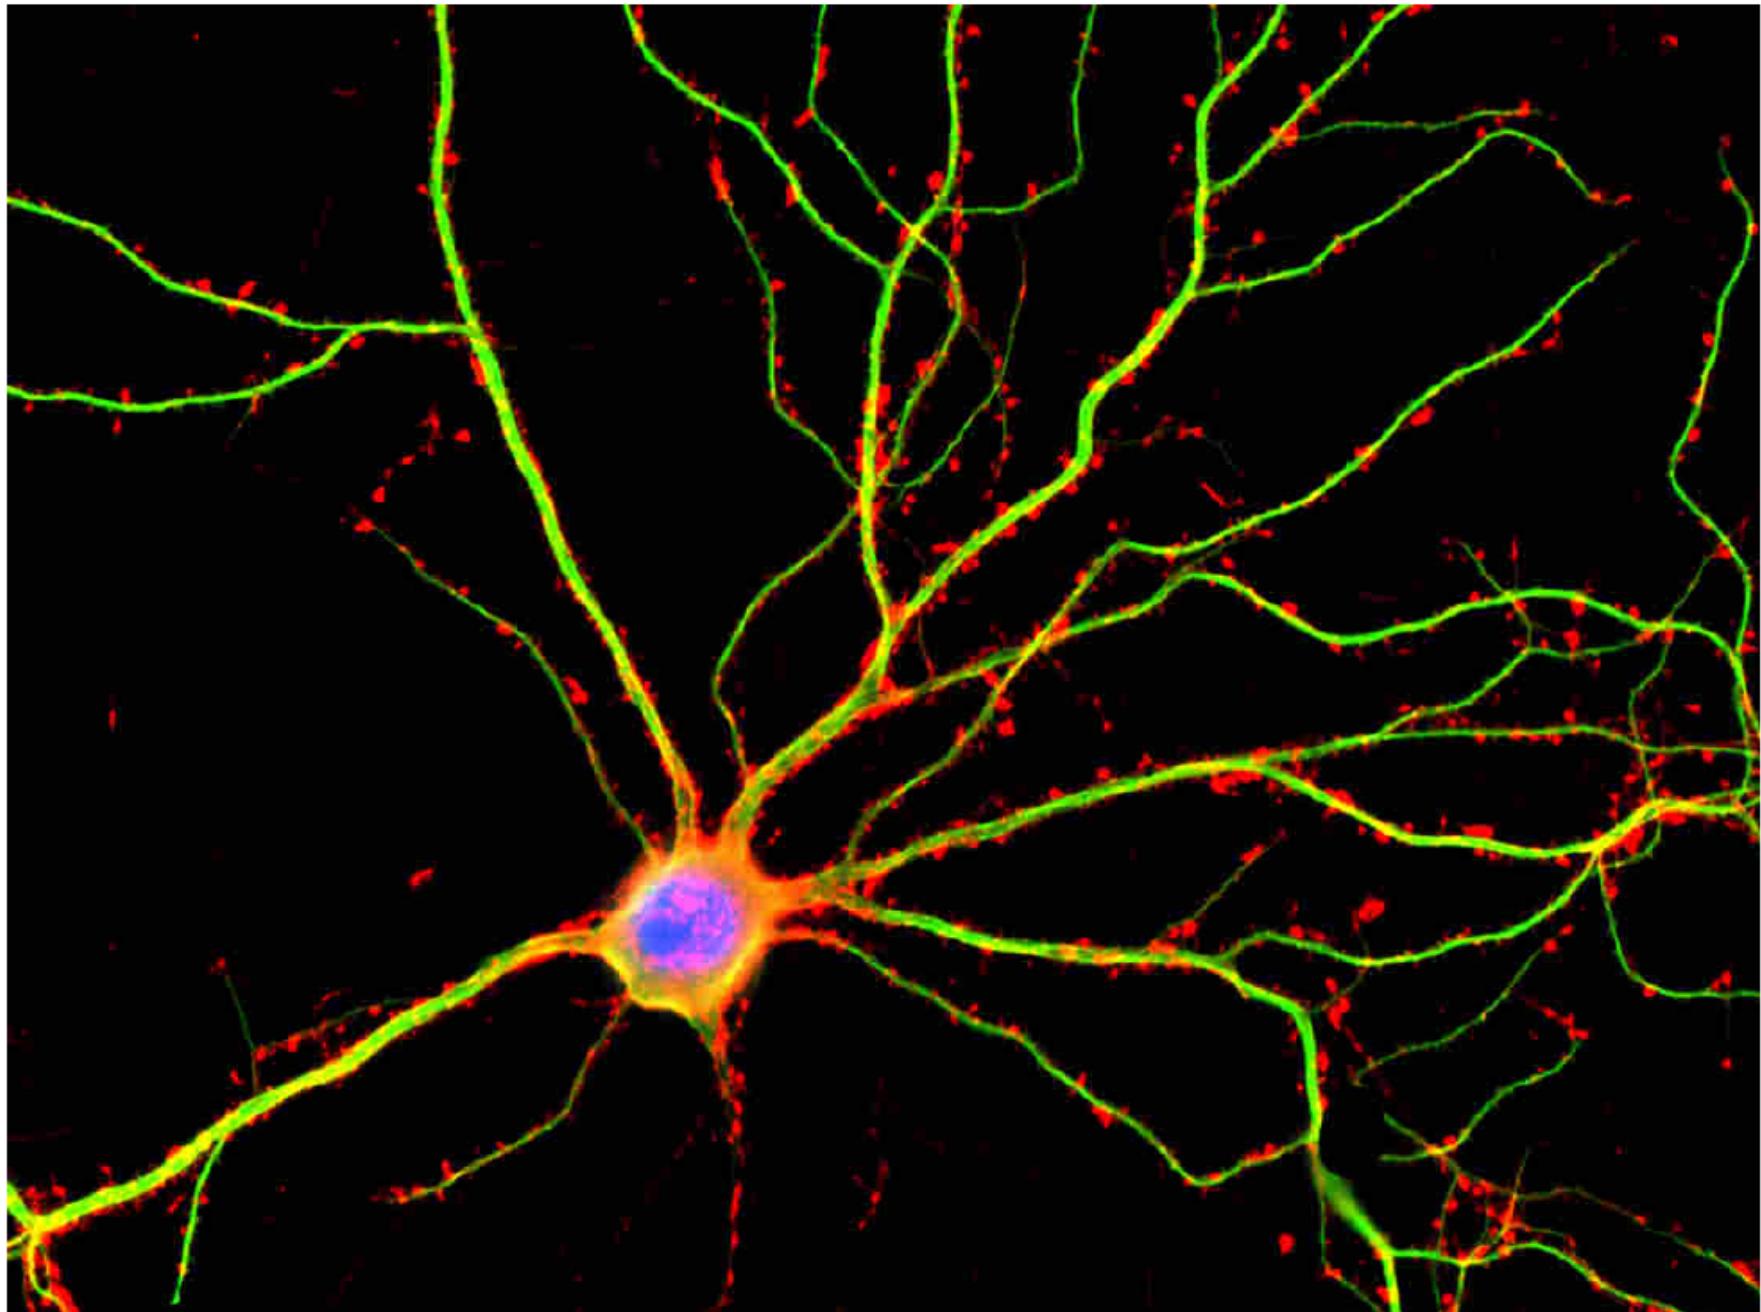
Training/
Determine
parameters

Basics Neural Networks

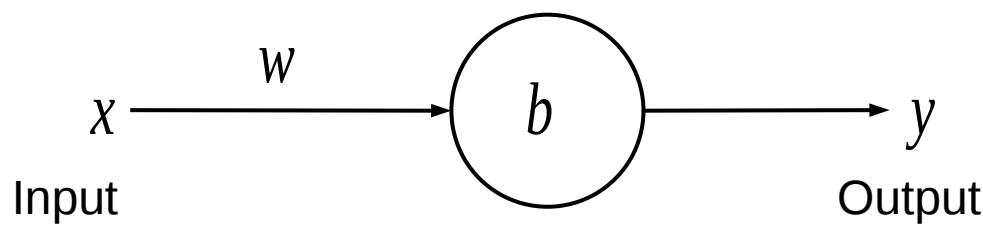
Basics Neuronal Networks

- What are artificial neural networks?
 - Biological motivated machine learning method
 - Consists out of artificial neurons (Perceptron)
- Classification
 - Supervised learning
 - Required training examples x_t and labels y_t
 - Learns a function $f: y = f(x)$
 - Symbolic and sub-symbolic ($y \in \mathbb{R}^N$)

Neuron

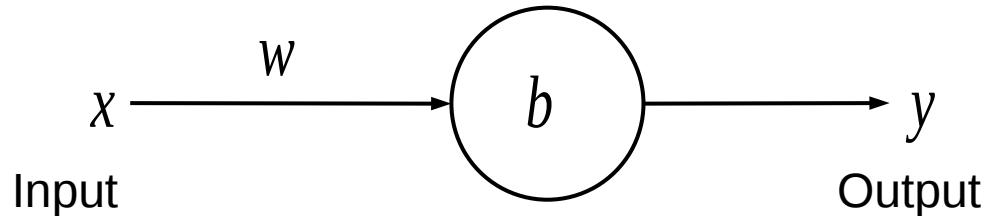


Perceptron (Neuron)

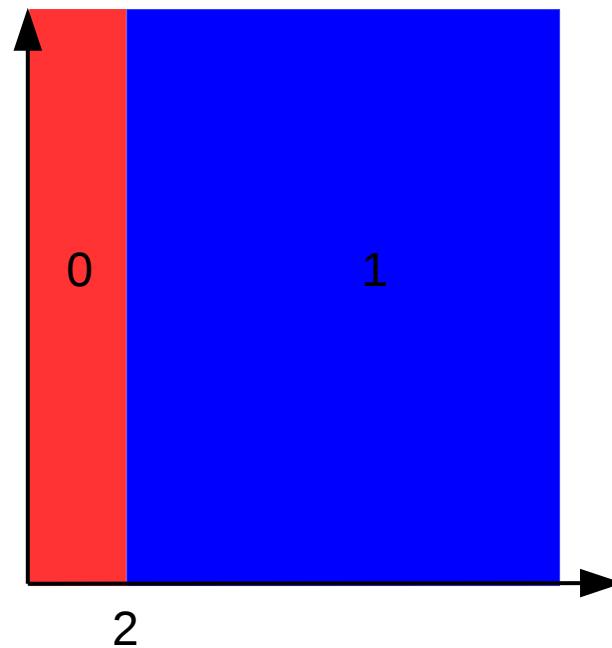


$$y = \begin{cases} 0 & \text{if } w \cdot x \leq b \\ 1 & \text{if } w \cdot x > b \end{cases}$$

Perceptron (Neuron) - Example

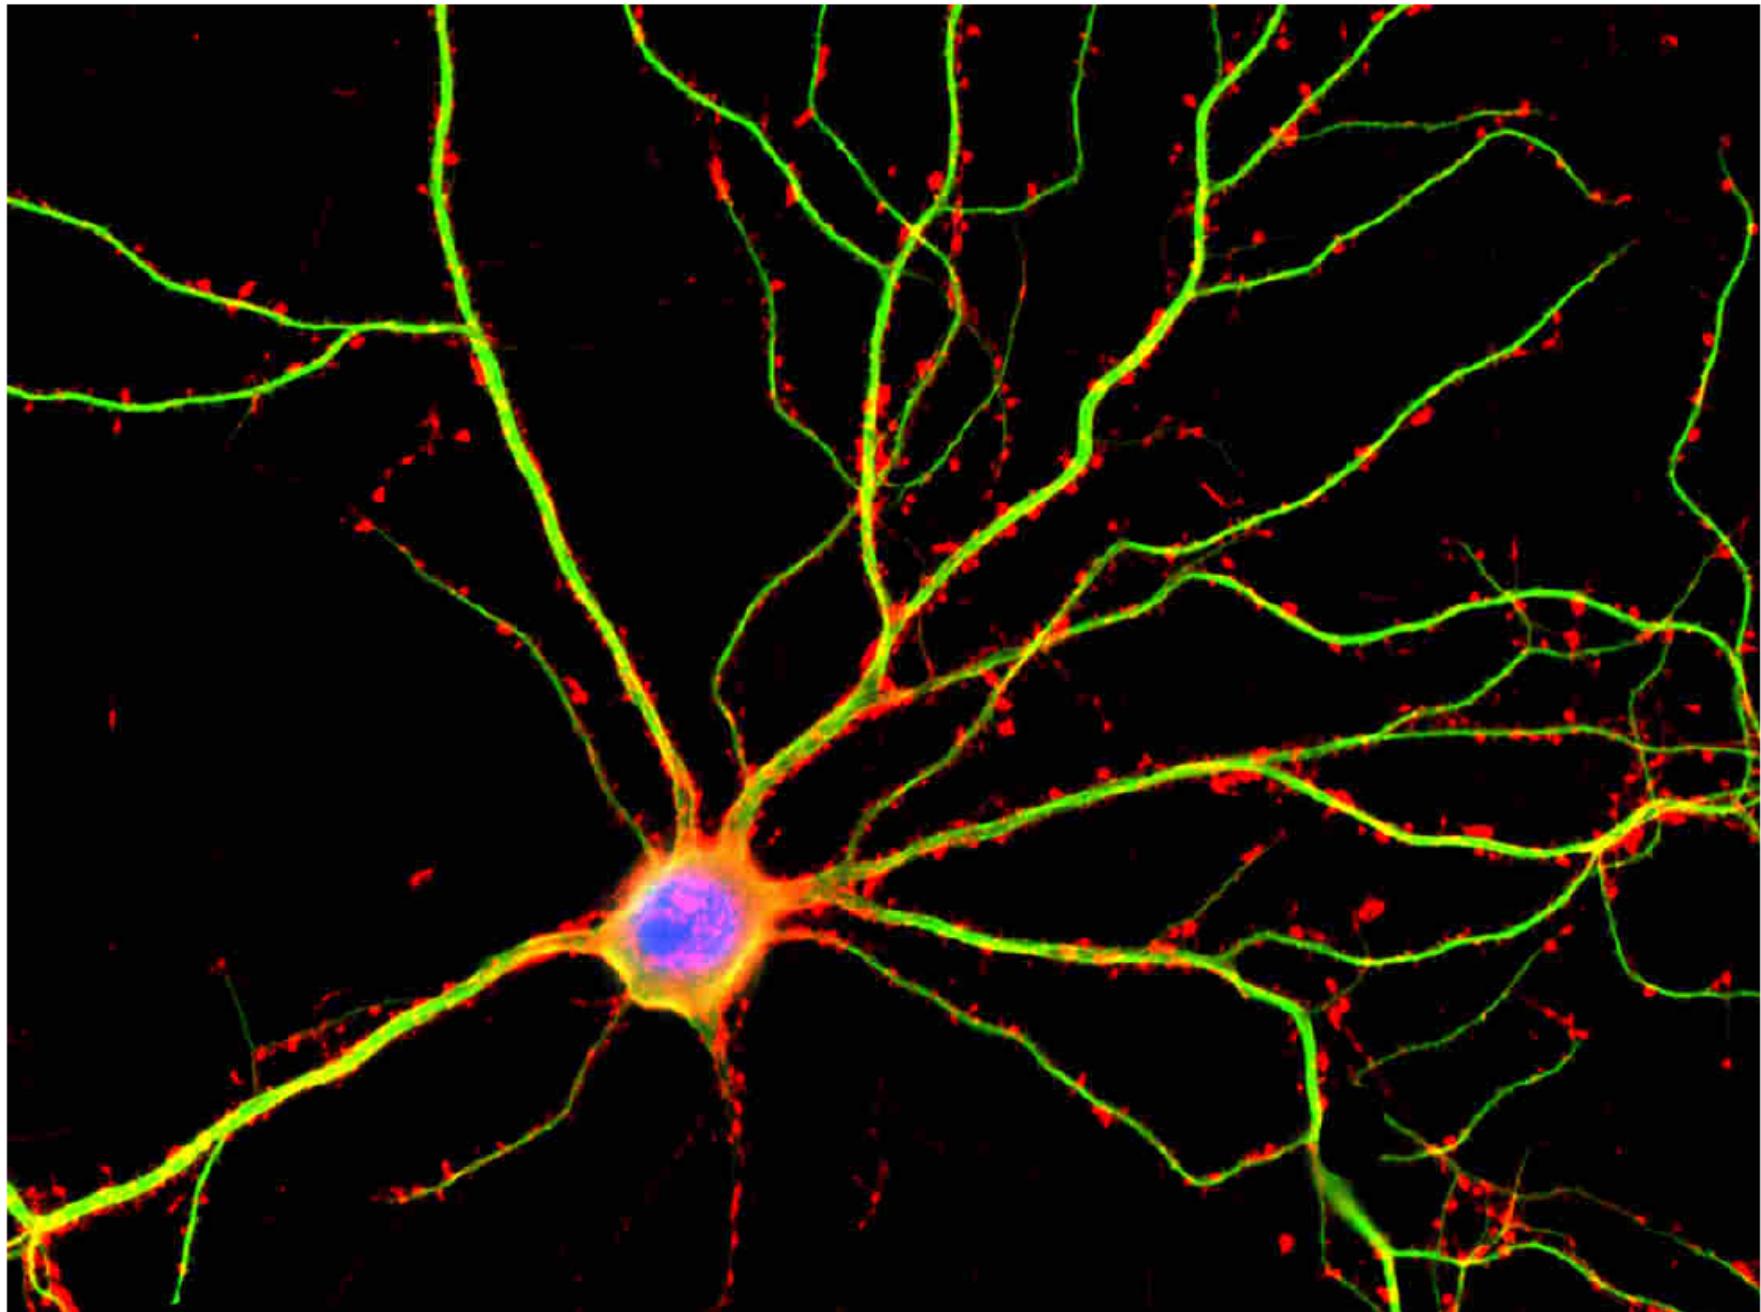


$$y = \begin{cases} 0 & \text{if } w \cdot x \leq b \\ 1 & \text{if } w \cdot x > b \end{cases}$$

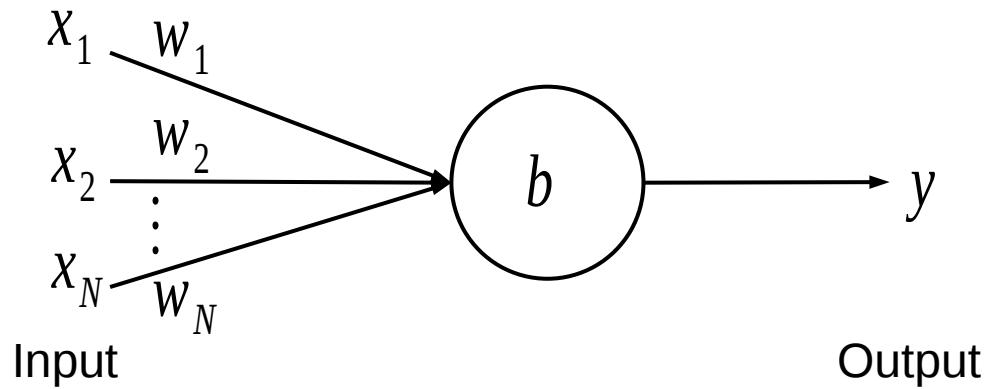


$$w = 1, b = 2$$

Neuron

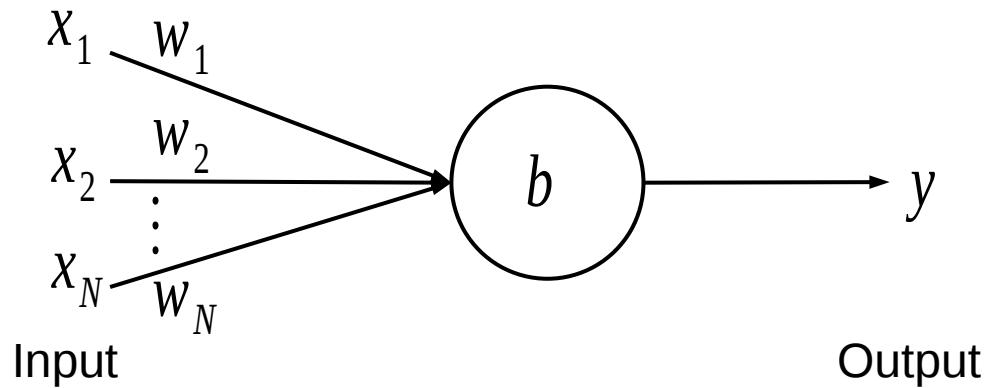


Perceptron (Neuron)



$$y = \begin{cases} 0 & \text{if } \sum_i w_i \cdot x_i \leq b \\ 1 & \text{if } \sum_i w_i \cdot x_i > b \end{cases}$$

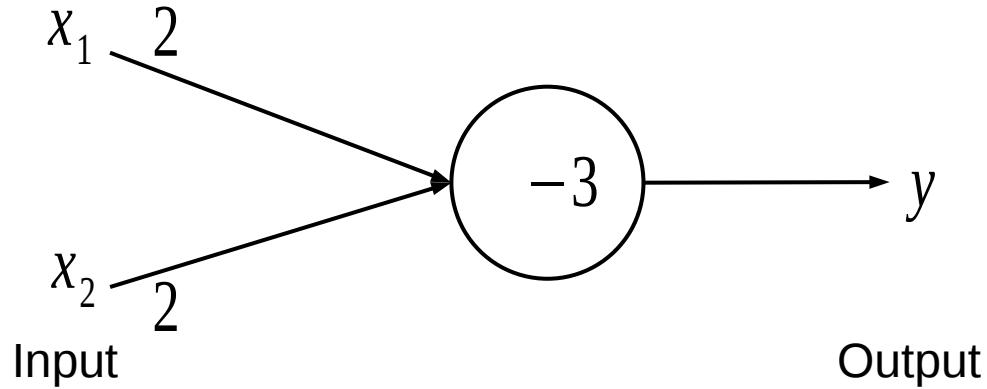
Perceptron (Neuron)



$$w \cdot x \equiv \sum_i w_i \cdot x_i$$
$$b \equiv -b$$

$$y = \begin{cases} 0 & \text{if } w \cdot x + b \leq 0 \\ 1 & \text{if } w \cdot x + b > 0 \end{cases}$$

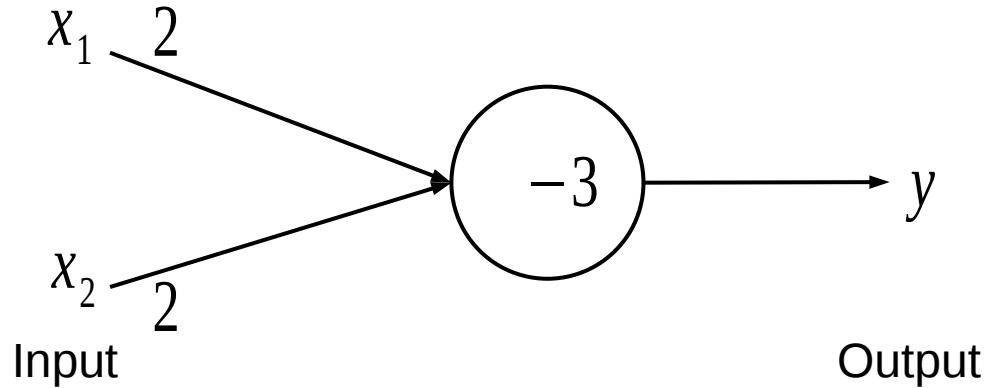
Perceptron (Neuron)



Assume: $x_1, x_2 \in \{0, 1\}$

- Which function is being computed?

Perceptron (Neuron)

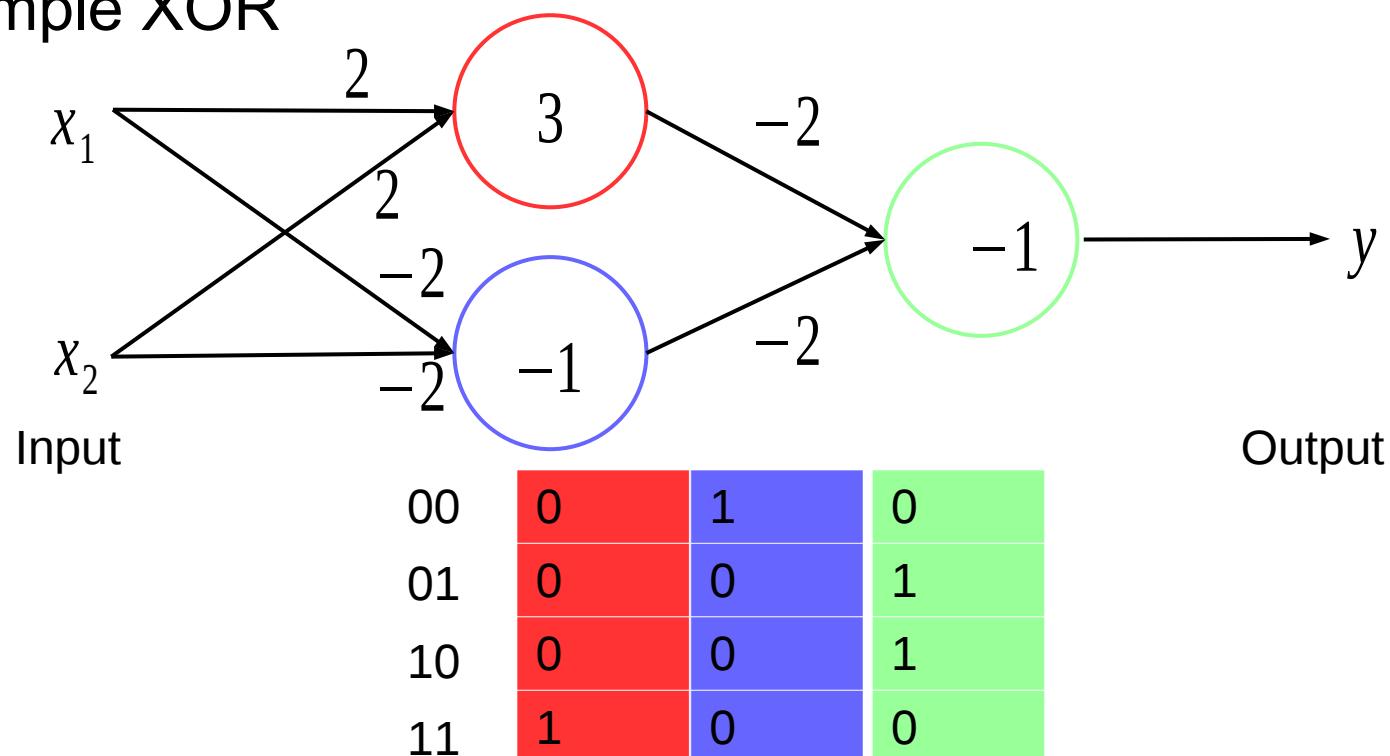


	$X_1 = 0$	$X_1 = 1$
$X_2 = 0$	$w \cdot x + b = -3$ $y = 0$	$w \cdot x = -1$ $y = 0$
$X_2 = 1$	$w \cdot x + b = -1$ $y = 0$	$w \cdot x = 1$ $y = 1$

=> Logical AND

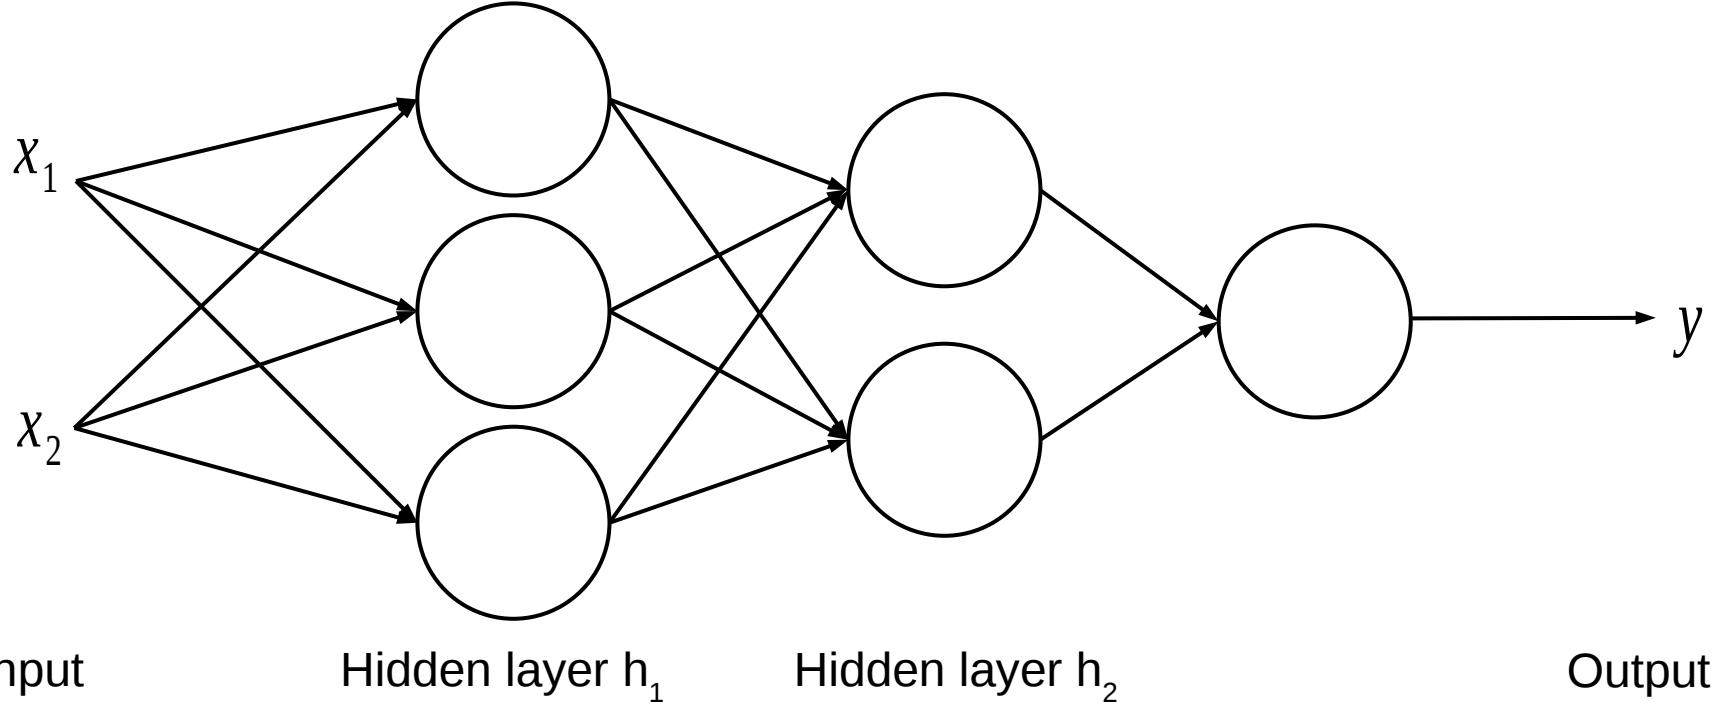
Multi-layered Network

- Perceptrons are linear discriminators
 - Usable for simple logical operations (AND, OR, NAND, ...)
 - Complex operations required more layers
- Example XOR



Multi-layered Network

- Multi-layered networks can represent arbitrary functions
 - Turing complete (assuming infinite depth)



Input

$$\begin{aligned} h_1 &= W_1 \cdot x + b_1 \\ h_2 &= W_2 \cdot h_1 + b_2 \\ y &= W_3 \cdot h_2 + b_3 \end{aligned}$$

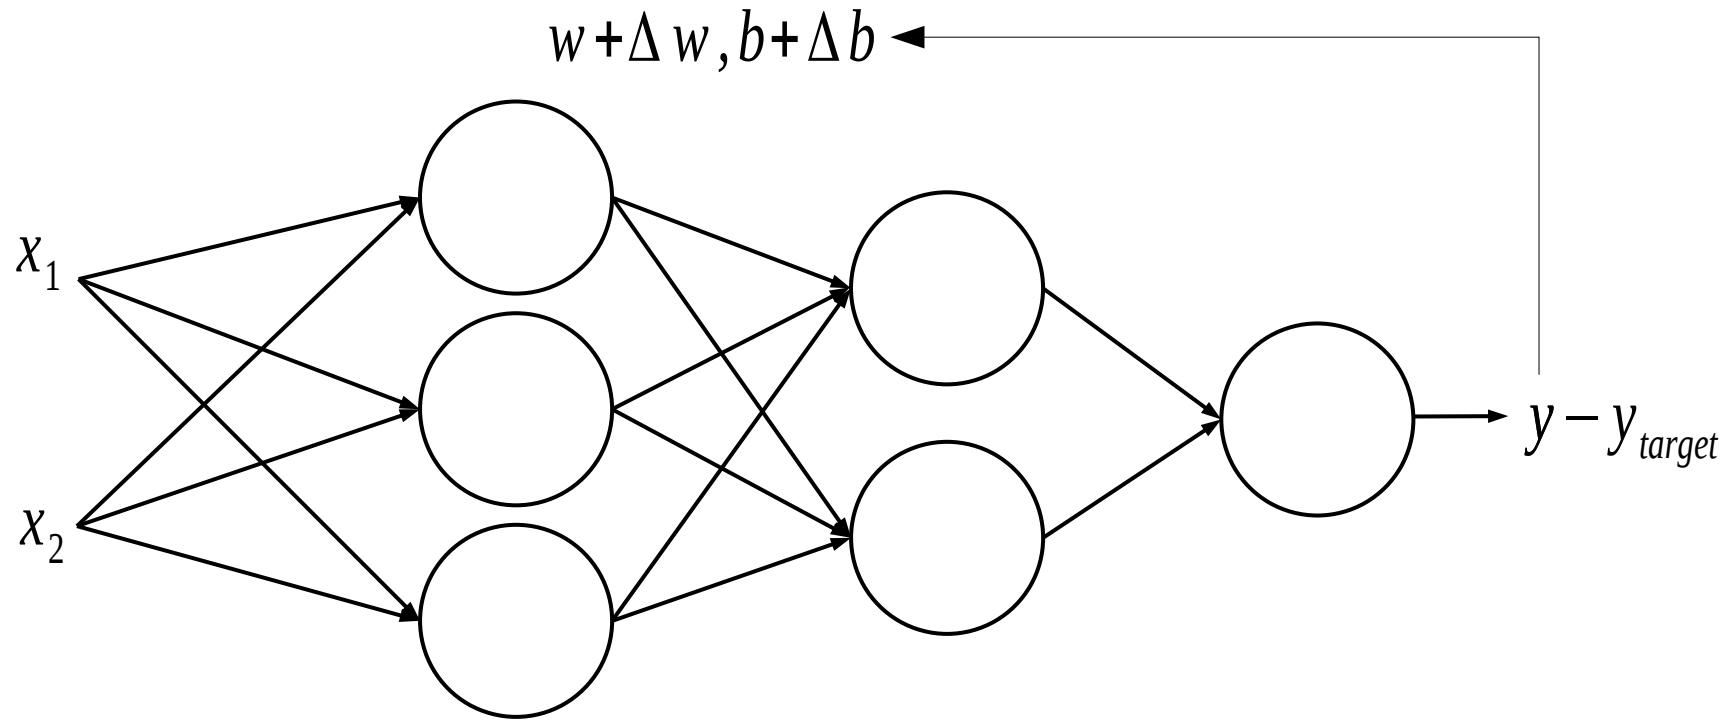
Hidden layer h_1

$$\begin{aligned} W_1 &= [w_{1,1} | w_{1,2} | w_{1,3}] \in \mathbb{R}^{2 \times 3}, b_1 = [b_{1,1} | b_{1,2} | b_{1,3}]^T \in \mathbb{R}^3 \\ W_2 &= [w_{2,1} | w_{2,2}] \in \mathbb{R}^{3 \times 2}, b_2 = [b_{2,1} | b_{2,2}]^T \in \mathbb{R}^2 \\ W_3 &= [w_{3,1}] \in \mathbb{R}^{2 \times 1}, b_3 = [b_{3,1}]^T \in \mathbb{R}^1 \end{aligned}$$

Output

Parameter estimation - Training

- Multi-layered networks can represent arbitrary functions
 - Given x and y
 - How to determine network parameters?

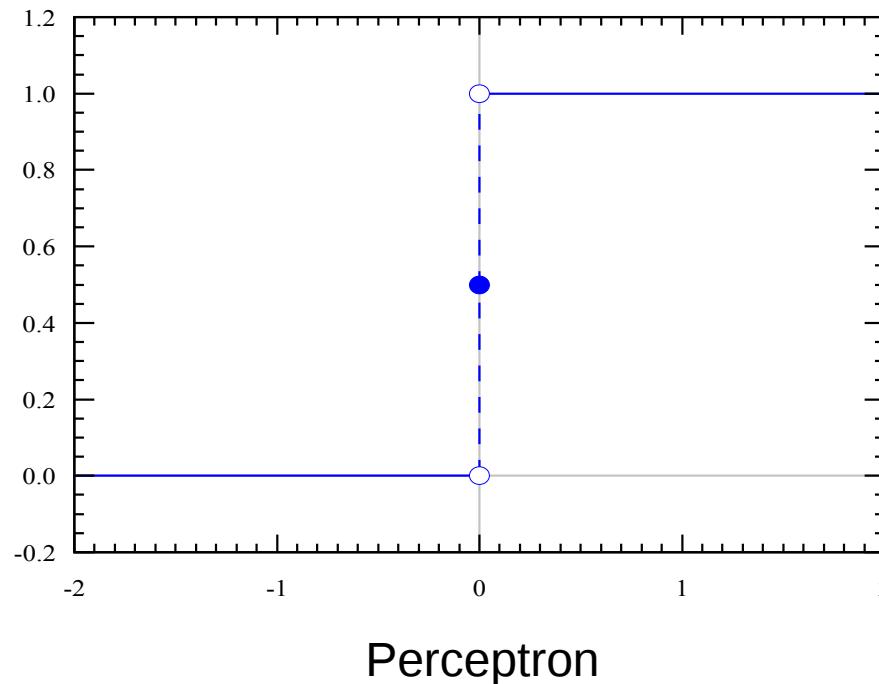


Input Hidden layer h_1 Hidden layer h_2 Output

- Through small changes of the weights, the output of the network approaches the solution: The network is learning

Training - Problem

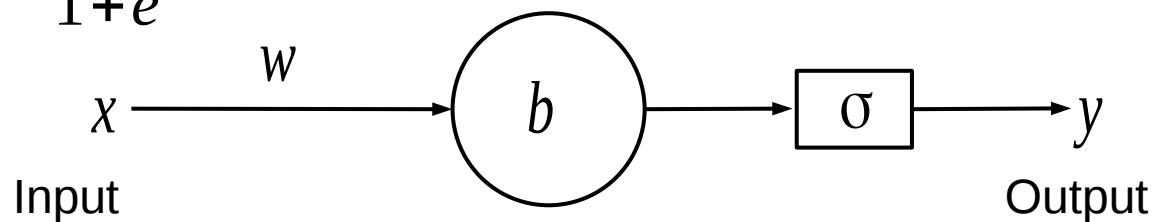
- But the network consists out of perceptrons
 - Small changes in weights can have large impact on the behavior of the entire network



Training - Problem

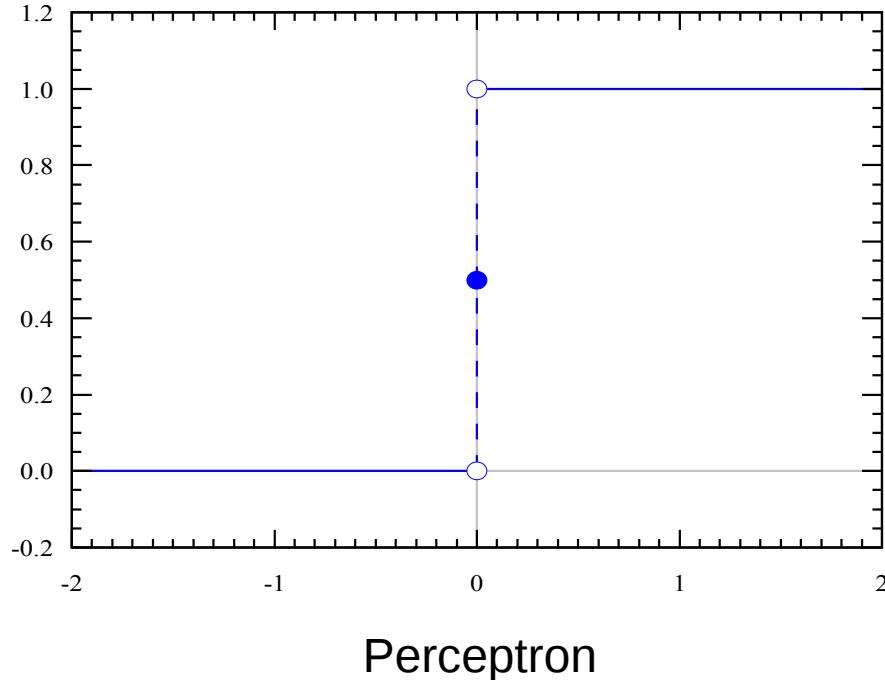
- But the network consists out of perceptrons
 - Small changes in weights can have large impact on the behavior of the entire network
- Solution: Sigmoid non-linearity

- $$\sigma(z) = \frac{1}{1+e^{-z}}$$

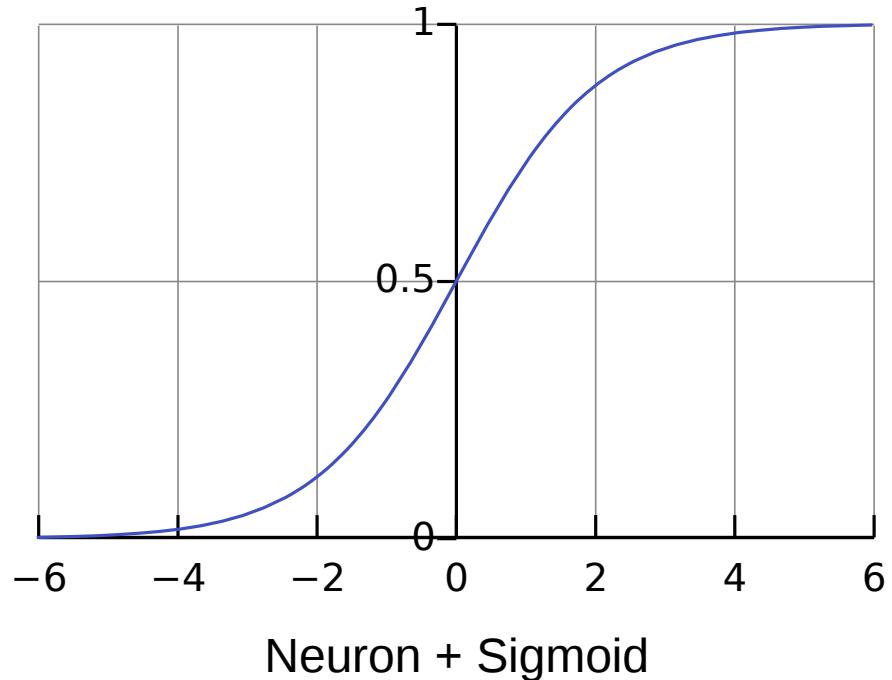


$$y = \sigma(w \cdot x + b)$$

Training - Sigmoid



Perceptron



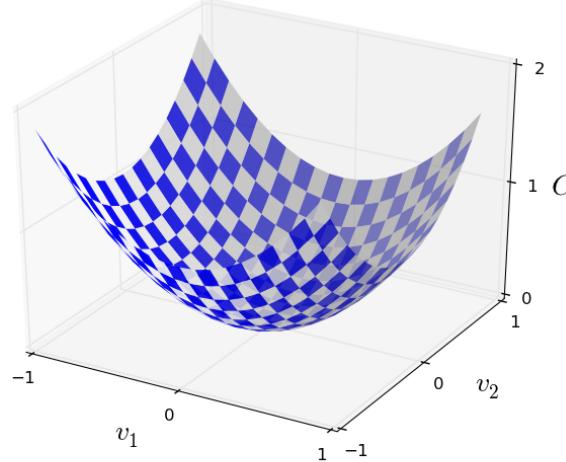
Neuron + Sigmoid

- Due to the smoothness, it is guaranteed that small changes in the weights only cause small changes in the output

Training – Gradient descent

- Optimization through minimization of cost function
 - Example: Euclidean norm

$$C(W, b) = \frac{1}{2n} \sum_{i=1}^n \|y(x_i) - y_{target}\|^2$$



- Gradient descent

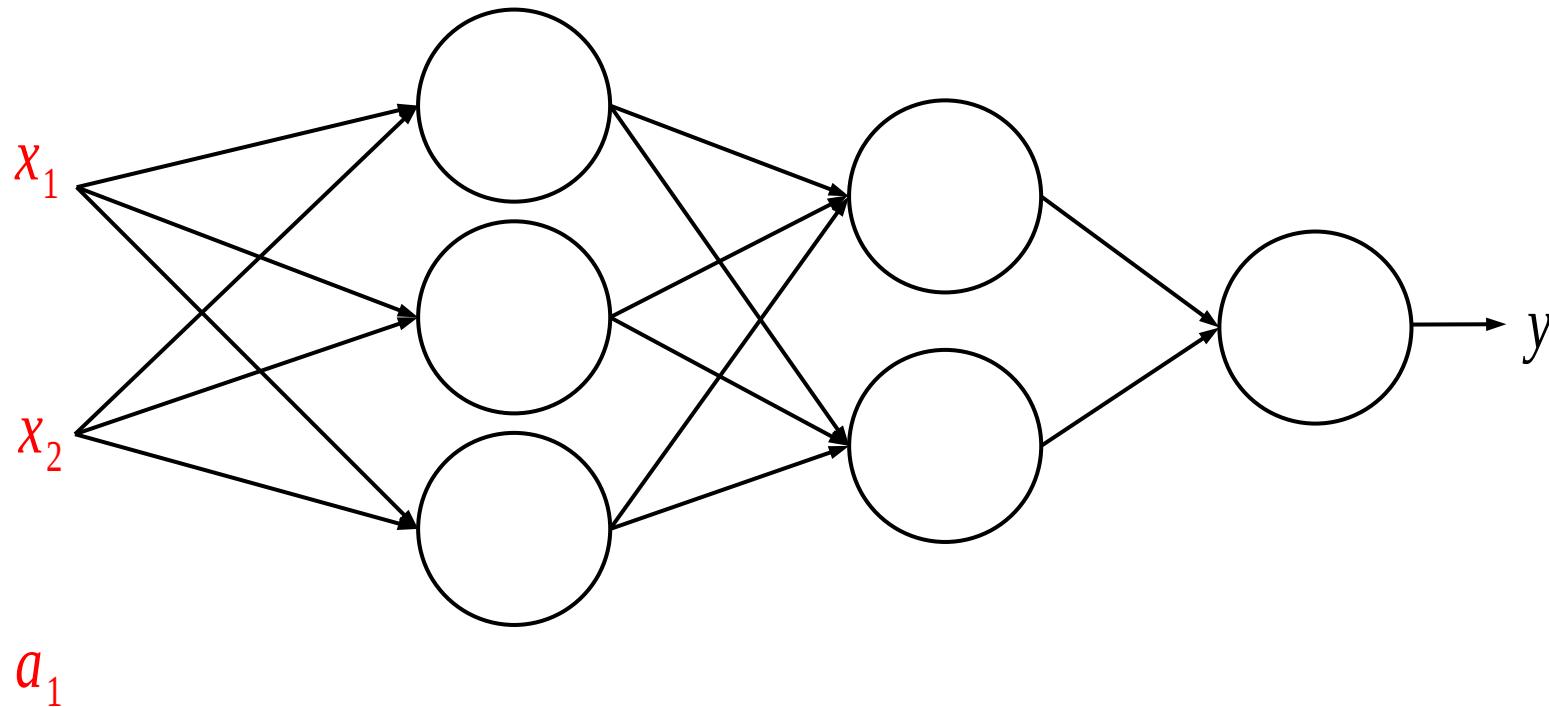
$$W_k = W_k - \eta \frac{\delta C}{\delta W_k}$$

$$b_k = b_k - \eta \frac{\delta C}{\delta b_k}$$

- η : Learning rate, determines step size
- η too large: Minimums are „skipped”
 - η too small: Local Minimums cannot be escaped

Training - Backpropagation

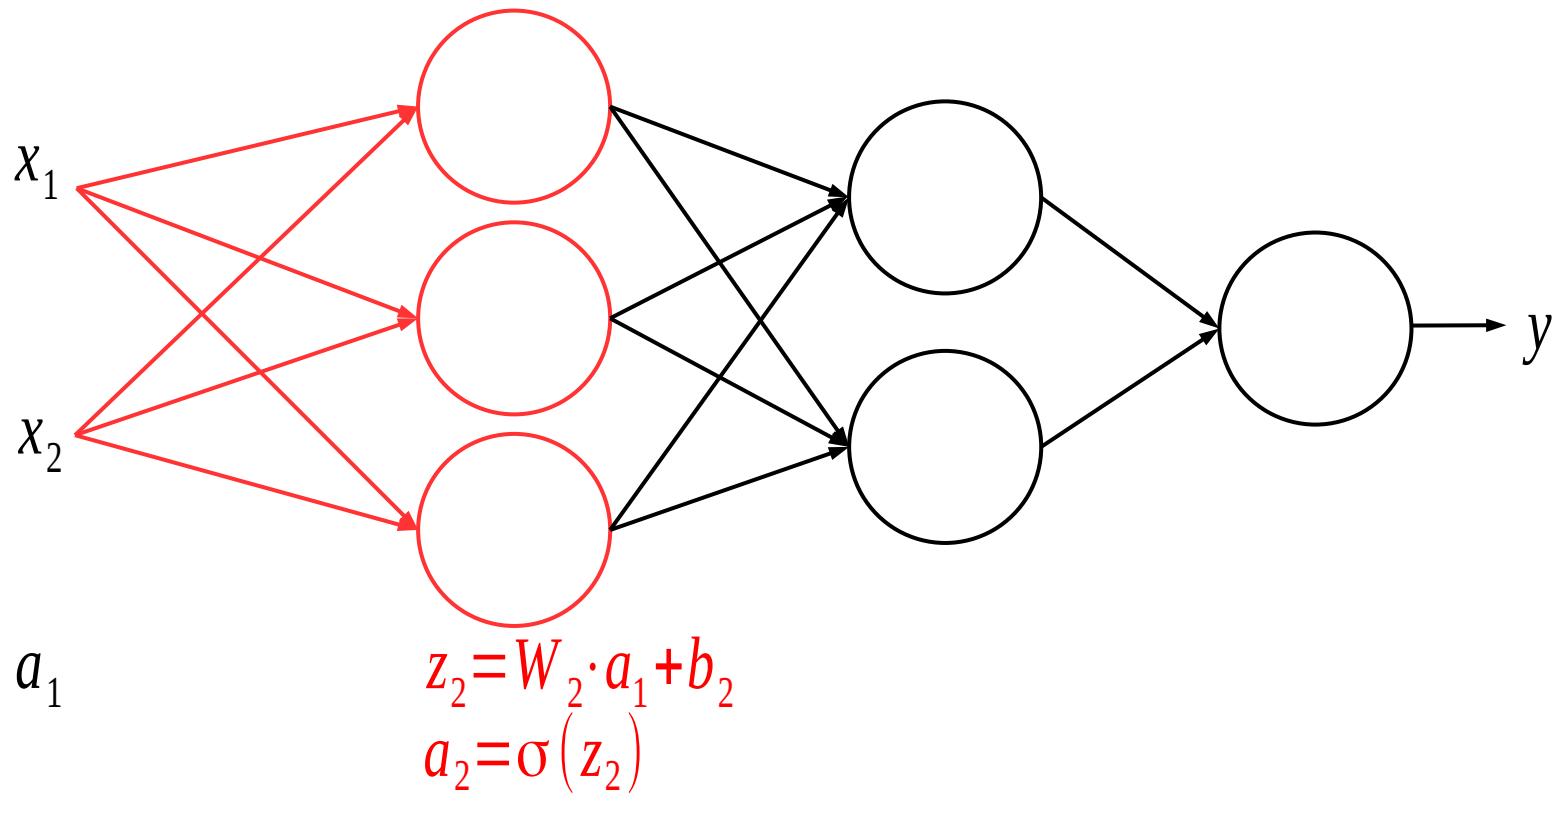
- Determine the partial derivatives through Backpropagation
- 1st step: Activation of the input layer $a_1 = x$



Training - Backpropagation

- 2nd step: For the remaining layers (2...L) compute:

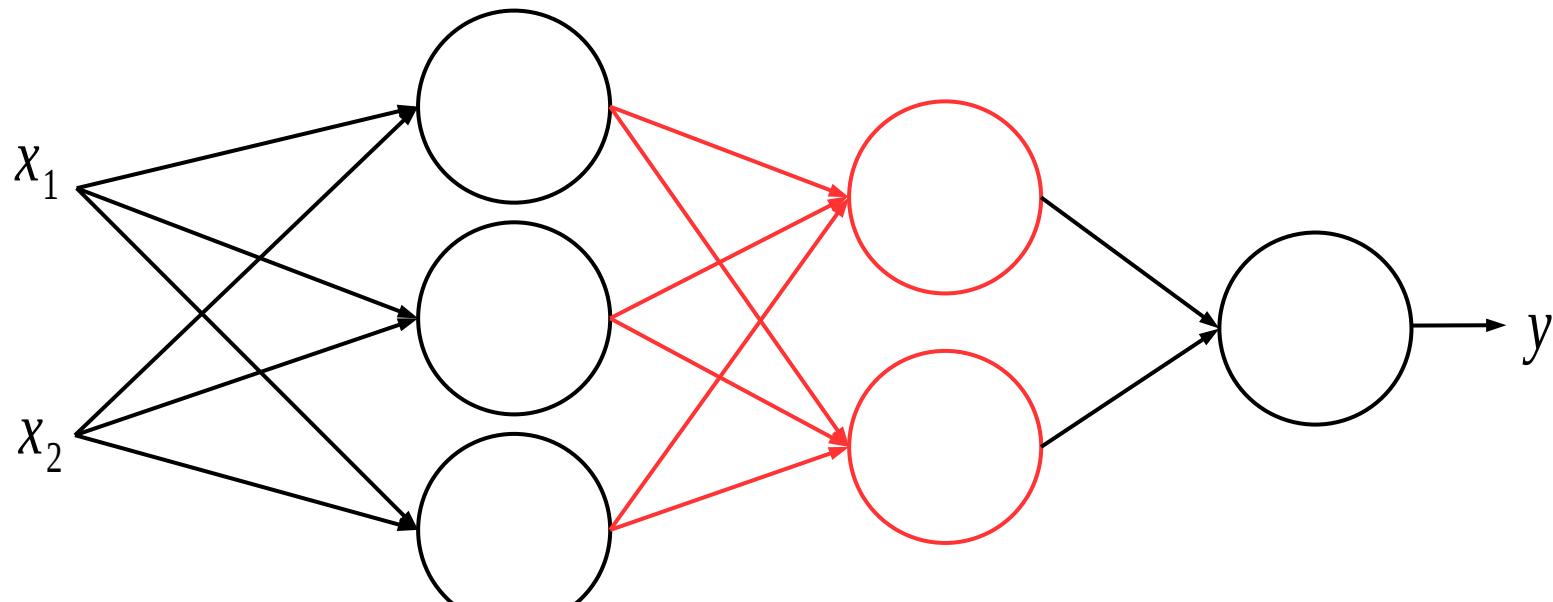
$$\begin{aligned} z_l &= W_l \cdot a_{l-1} + b_l \\ a_l &= \sigma(z_l) \end{aligned} \quad l = 2, \dots, L$$



Training - Backpropagation

- 2nd step: For the remaining layers ($2 \dots L$) compute:

$$\begin{aligned} z_l &= W_l \cdot a_{l-1} + b_l \\ a_l &= \sigma(z_l) \end{aligned} \quad l = 2, \dots, L$$



a_1

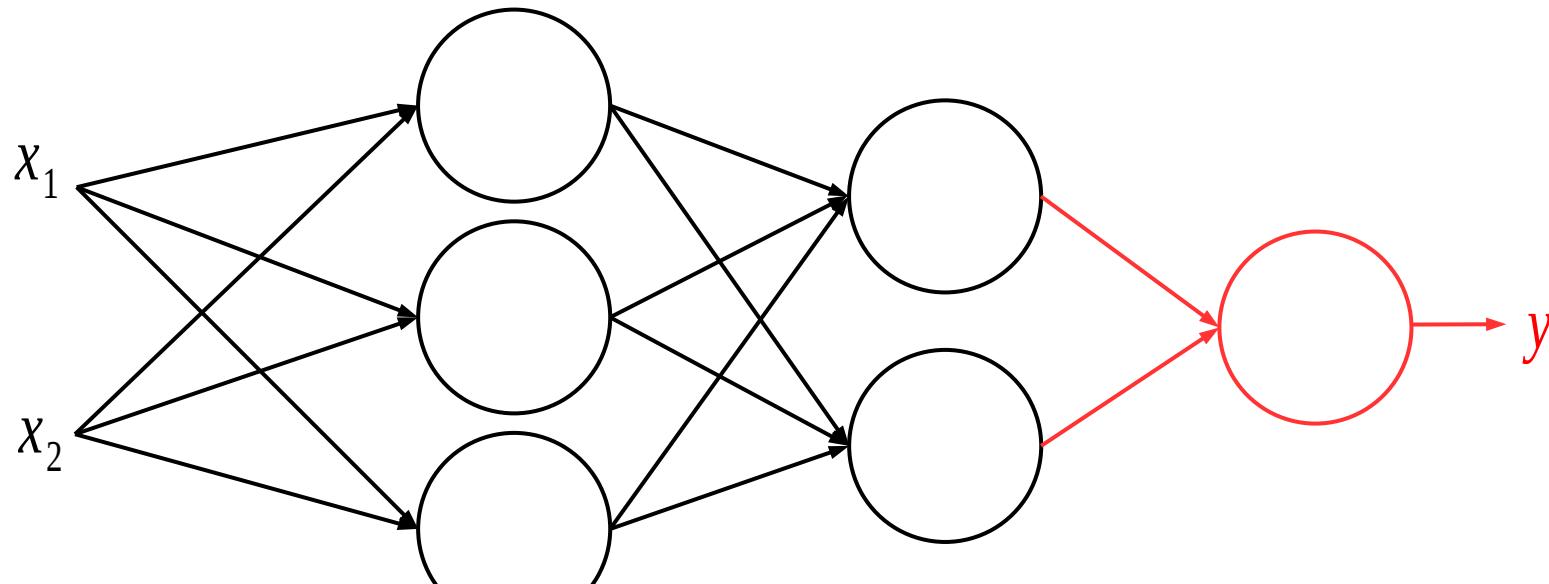
$$\begin{aligned} z_2 &= W_2 \cdot a_1 + b_2 \\ a_2 &= \sigma(z_2) \end{aligned}$$

$$\begin{aligned} z_3 &= W_3 \cdot a_2 + b_3 \\ a_3 &= \sigma(z_3) \end{aligned}$$

Training - Backpropagation

- 2nd step: For the remaining layers (2...L) compute:

$$\begin{aligned} z_l &= W_l \cdot a_{l-1} + b_l \\ a_l &= \sigma(z_l) \end{aligned} \quad l = 2, \dots, L$$



$$\begin{aligned} a_1 \\ z_2 &= W_2 \cdot a_1 + b_2 \\ a_2 &= \sigma(z_2) \end{aligned}$$

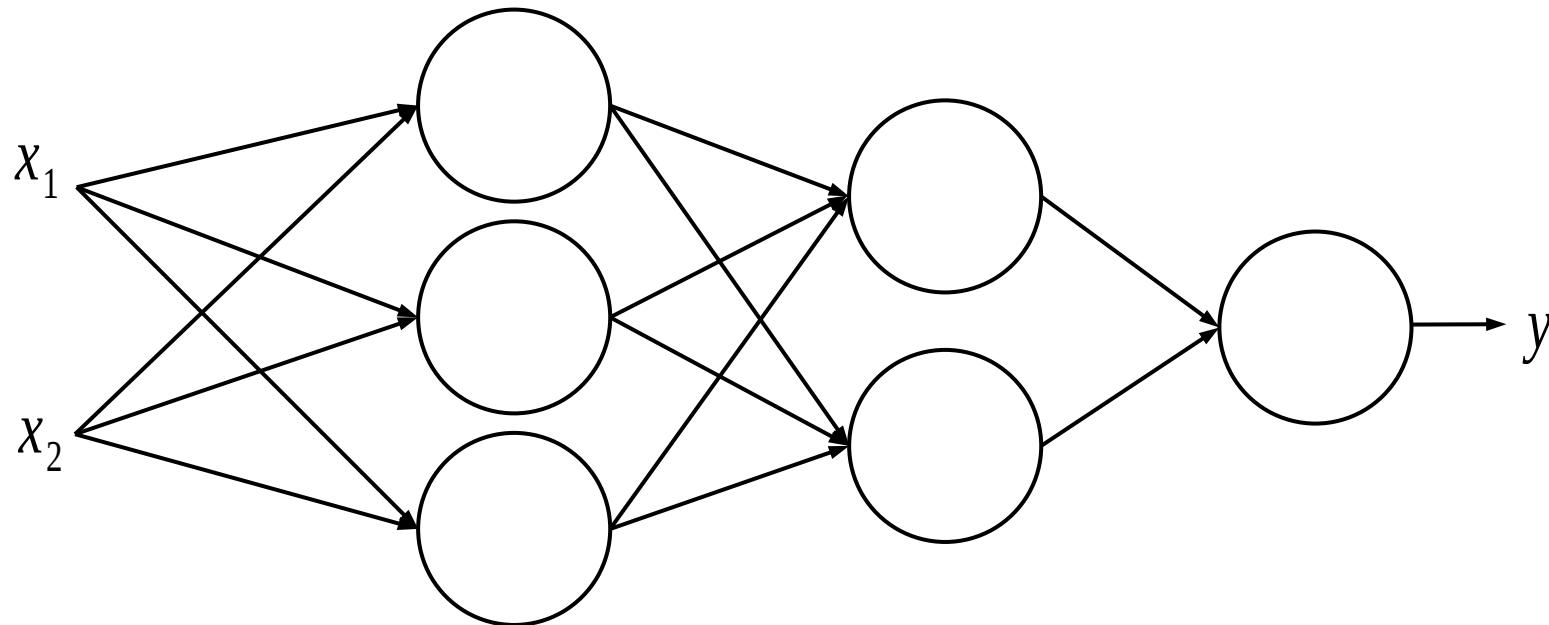
$$\begin{aligned} z_3 &= W_3 \cdot a_2 + b_3 \\ a_3 &= \sigma(z_3) \end{aligned}$$

$$\begin{aligned} z_4 &= W_4 \cdot a_3 + b_4 \\ y &= a_4 = \sigma(z_4) \end{aligned}$$

Training - Backpropagation

- 3rd step: Compute output error

$$C(W, b) = \frac{1}{2n} \sum_{i=1}^n \|a_L - y_{target}\|^2$$

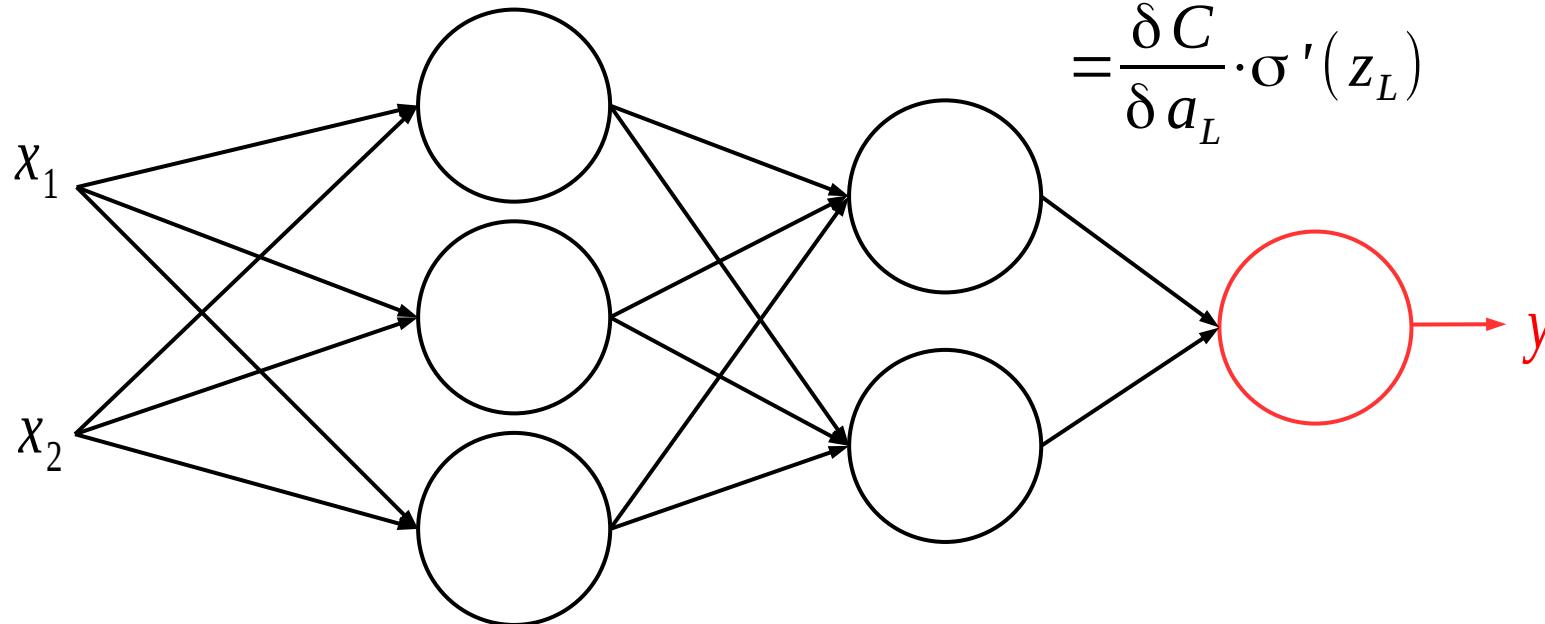


Training - Backpropagation

- 3rd step: Compute output error

$$C(W, b) = \frac{1}{2n} \sum_{i=1}^n \|a_L - y_{target}\|^2$$

$$\begin{aligned}\delta_L &= \frac{\delta C}{\delta z_L} \\ &= \frac{\delta C}{\delta a_L} \cdot \frac{\delta a_L}{\delta z_L} \\ &= \frac{\delta C}{\delta a_L} \cdot \sigma'(z_L)\end{aligned}$$



$$a_1$$

$$\begin{aligned}z_2 &= W_2 \cdot a_1 + b_2 \\ a_2 &= \sigma(z_2)\end{aligned}$$

$$\begin{aligned}z_3 &= W_3 \cdot a_2 + b_3 \\ a_3 &= \sigma(z_3)\end{aligned}$$

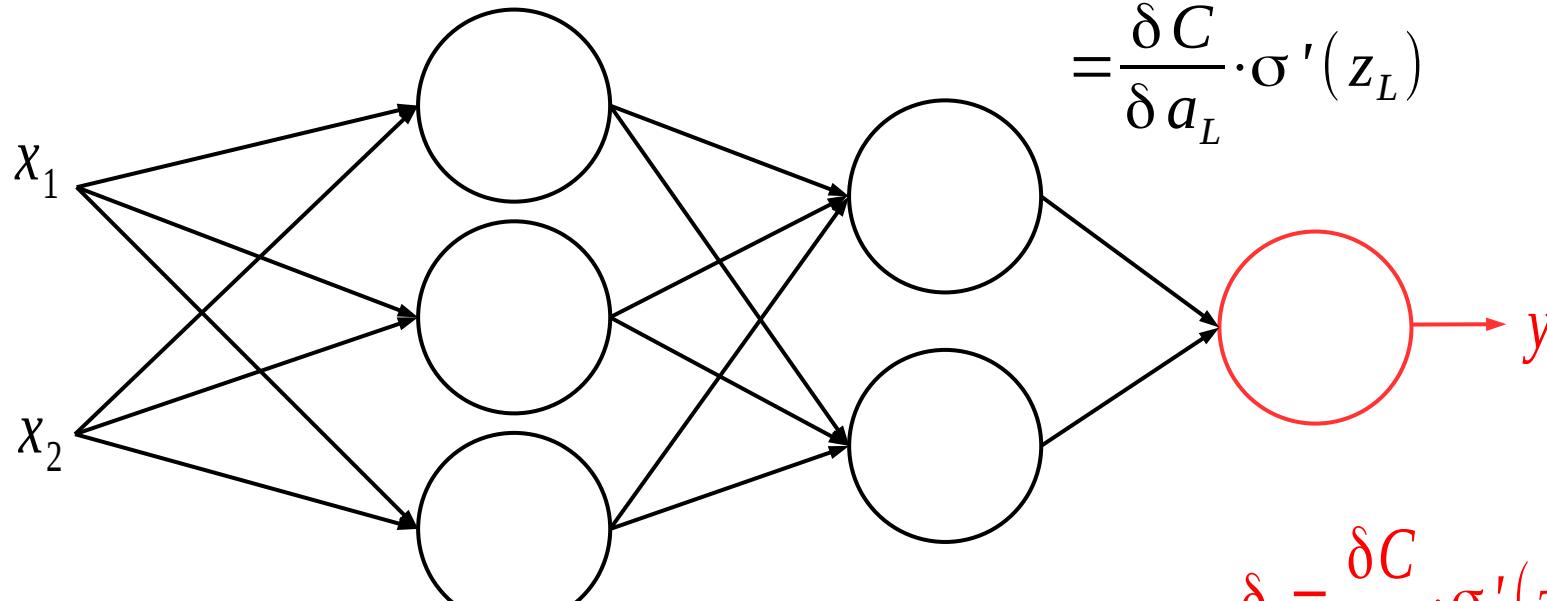
$$\begin{aligned}z_4 &= W_4 \cdot a_3 + b_4 \\ y &= a_4 = \sigma(z_4)\end{aligned}$$

Training - Backpropagation

- 3rd step: Compute output error

$$C(W, b) = \frac{1}{2n} \sum_{i=1}^n \|a_L - y_{target}\|^2$$

$$\begin{aligned}\delta_L &= \frac{\delta C}{\delta z_L} \\ &= \frac{\delta C}{\delta a_L} \cdot \frac{\delta a_L}{\delta z_L} \\ &= \frac{\delta C}{\delta a_L} \cdot \sigma'(z_L)\end{aligned}$$



$$\delta_4 = \frac{\delta C}{\delta a_4} \cdot \sigma'(z_4)$$

a_1

$$\begin{aligned}z_2 &= W_2 \cdot a_1 + b_2 \\ a_2 &= \sigma(z_2)\end{aligned}$$

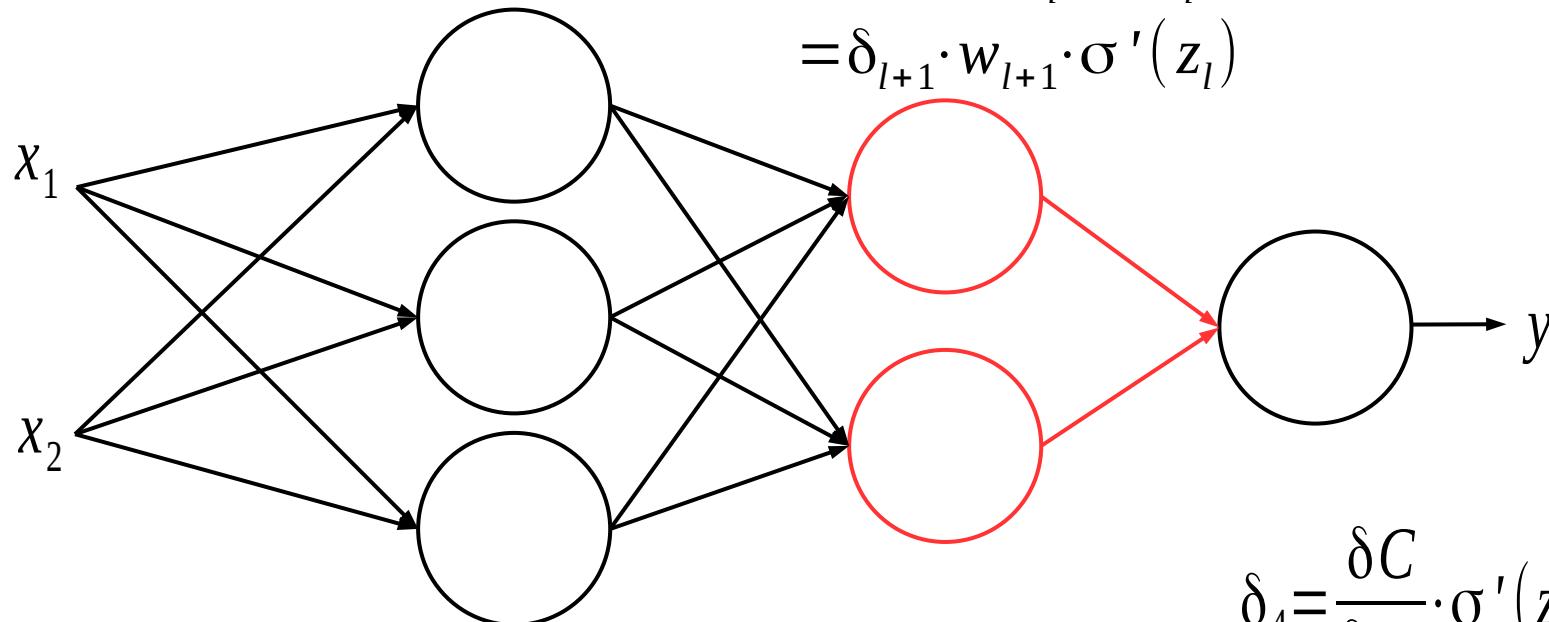
$$\begin{aligned}z_3 &= W_3 \cdot a_2 + b_3 \\ a_3 &= \sigma(z_3)\end{aligned}$$

$$\begin{aligned}z_4 &= W_4 \cdot a_3 + b_4 \\ a_4 &= \sigma(z_4)\end{aligned}$$

Training - Backpropagation

- 4th step: Propage output error back to the remaining layers

$$\begin{aligned}\delta_l &= \frac{\delta C}{\delta z_l} = \frac{\delta C}{\delta z_{l+1}} \cdot \frac{\delta z_{l+1}}{\delta z_l} \\ &= \delta_{l+1} \cdot \frac{\delta z_{l+1}}{\delta a_l} \cdot \frac{\delta a_l}{\delta z_l} \\ &= \delta_{l+1} \cdot w_{l+1} \cdot \sigma'(z_l)\end{aligned}\quad l=L-1, \dots, 2$$

 a_1

$$\begin{aligned}z_2 &= W_2 \cdot a_1 + b_2 \\ a_2 &= \sigma(z_2)\end{aligned}$$

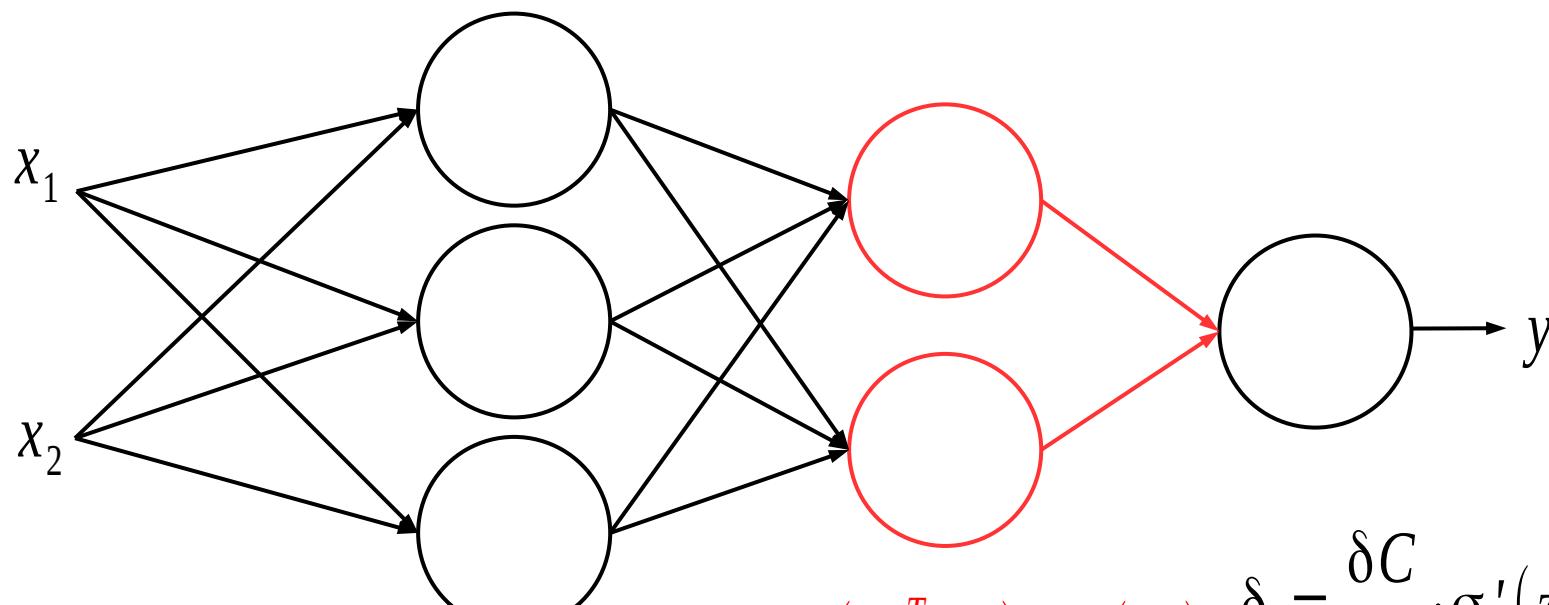
$$\begin{aligned}z_3 &= W_3 \cdot a_2 + b_3 \\ a_3 &= \sigma(z_3)\end{aligned}$$

$$\begin{aligned}z_4 &= W_4 \cdot a_3 + b_4 \\ y &= a_4 = \sigma(z_4)\end{aligned}$$

Training - Backpropagation

- 4th step: Propage output error back to the remaining layers

$$\delta_l = (W_{l+1}^T \cdot \delta_{l+1}) \circ \sigma'(z_l) \quad l = L-1, \dots, 2$$



$$\delta_3 = (W_4^T \cdot \delta_4) \circ \sigma'(z_4) \quad \delta_4 = \frac{\delta C}{\delta a_4} \cdot \sigma'(z_4)$$

a_1

$$z_2 = W_2 \cdot a_1 + b_2 \\ a_2 = \sigma(z_2)$$

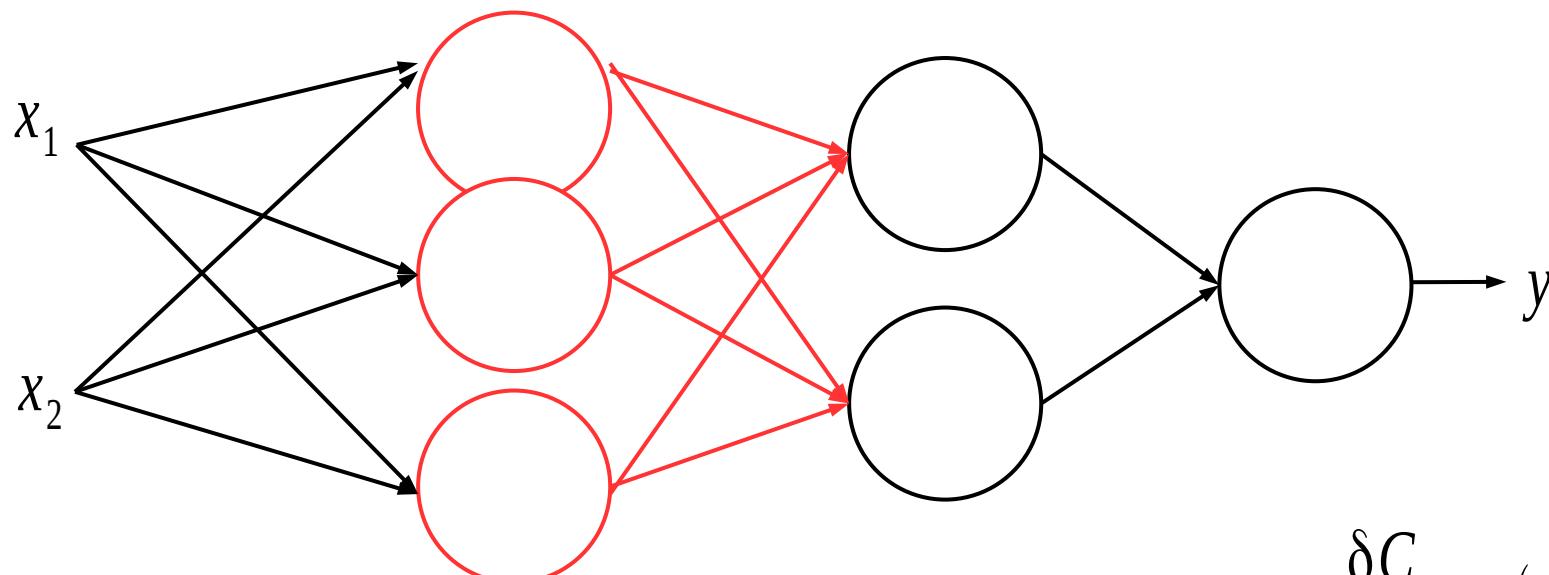
$$z_3 = W_3 \cdot a_2 + b_3 \\ a_3 = \sigma(z_3)$$

$$z_4 = W_4 \cdot a_3 + b_4 \\ a_4 = \sigma(z_4)$$

Training - Backpropagation

- 4th step: Propage output error back to the remaining layers

$$\delta_l = (W_{l+1}^T \cdot \delta_{l+1}) \circ \sigma'(z_l) \quad l = L-1, \dots, 2$$



$$\delta_2 = (W_3^T \cdot \delta_3) \circ \sigma'(z_2) \quad \delta_3 = (W_4^T \cdot \delta_4) \circ \sigma'(z_3) \quad \delta_4 = \frac{\delta C}{\delta a_4} \cdot \sigma'(z_4)$$

a_1

$$z_2 = W_2 \cdot a_1 + b_2 \\ a_2 = \sigma(z_2)$$

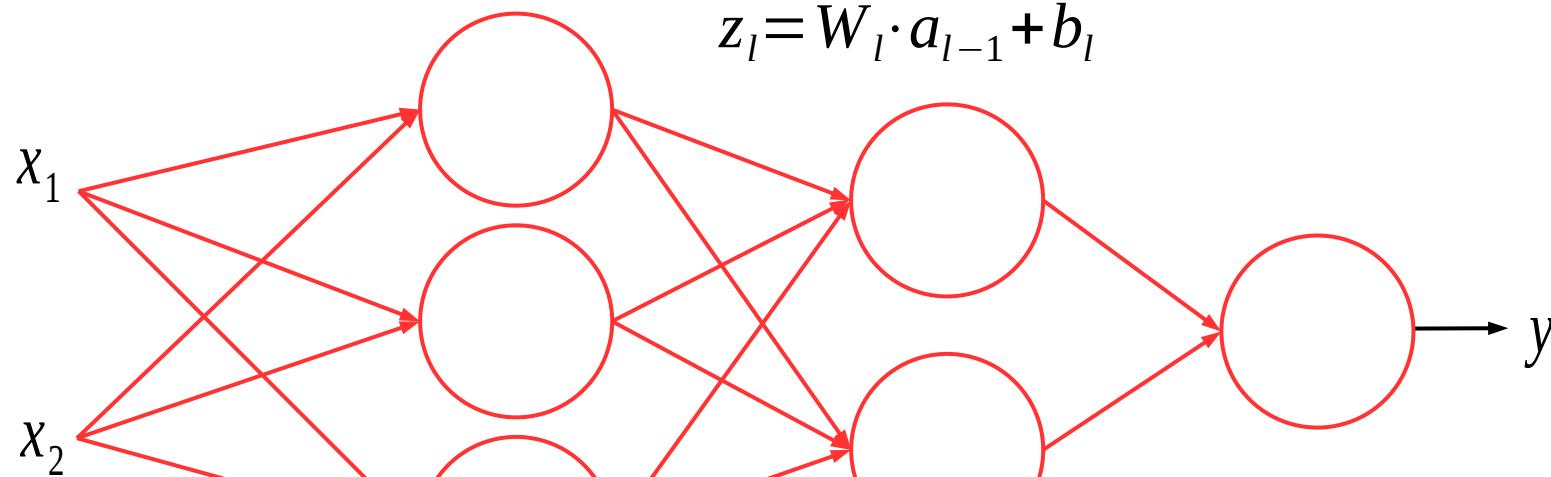
$$z_3 = W_3 \cdot a_2 + b_3 \\ a_3 = \sigma(z_3)$$

$$z_4 = W_4 \cdot a_3 + b_4 \\ a_4 = \sigma(z_4)$$

Training - Backpropagation

- 5th step: Compute partial derivatives

$$\frac{\delta C}{\delta W_l} = \frac{\delta C}{\delta z_l} \cdot \frac{\delta z_l}{\delta W_l} = \frac{\delta z_l}{\delta W_l} \cdot \delta_l^T = a_{l-1} \cdot \delta_l^T \quad \frac{\delta C}{\delta b_l} = \frac{\delta C}{\delta z_l} \cdot \frac{\delta z_l}{\delta b_l} = \delta_l$$

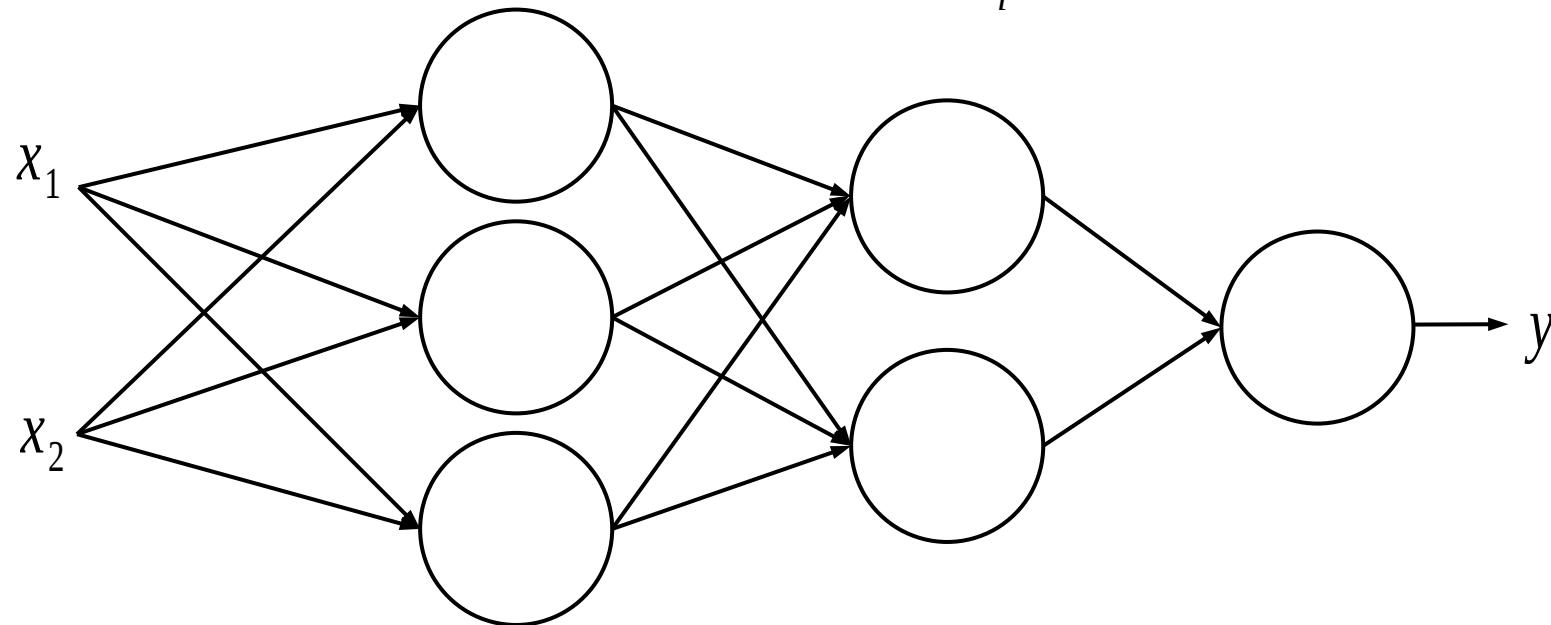


Training - Backpropagation

- 6th step: Gradient descent

$$W_l = W_l - \eta \frac{\delta C}{\delta W_l}$$

$$b_l = b_l - \eta \frac{\delta C}{\delta b_l}$$

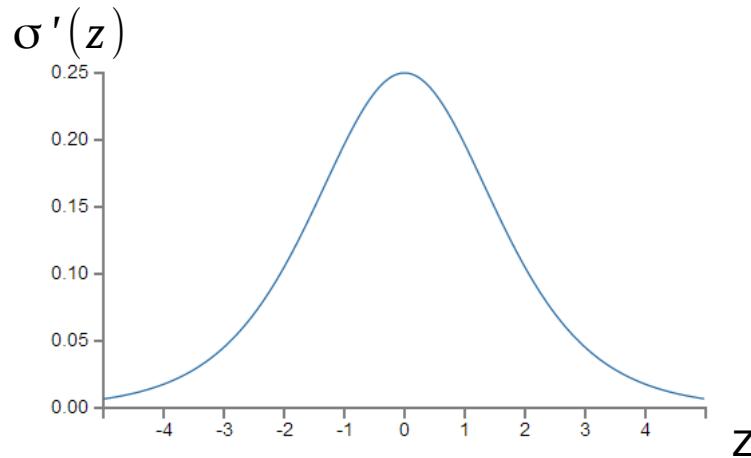


Deep networks: Deep Learning

- Arbitrary deep networks are powerful (Turing-complete)
 - Problem when training with Backpropagation

$$\delta_L = \frac{\delta C}{\delta a_L} \cdot \sigma'(z_L)$$

$$\delta_l = (w_{l+1}^T \cdot \delta_{l+1}) \circ \sigma'(z_l) \quad l=L-1, \dots, 2$$



- Through the influence of the derivative of the sigmoid, the gradient is shrinking

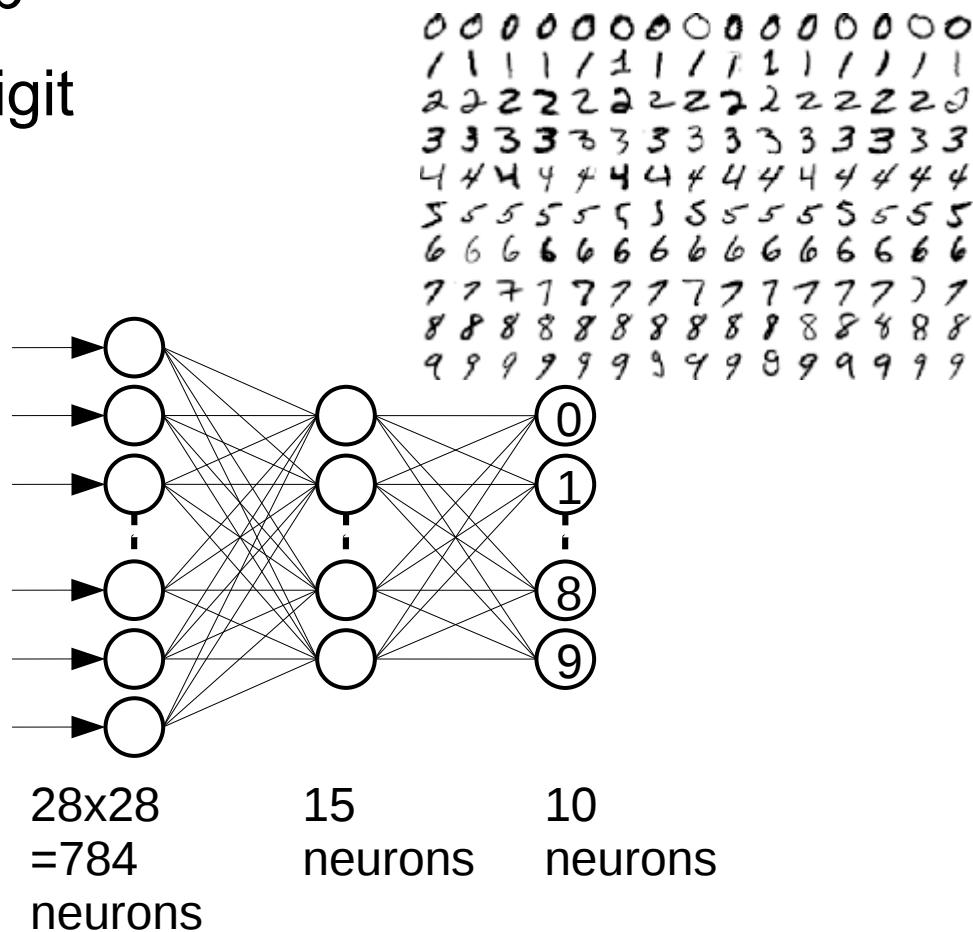
=> Vanishing Gradients

Deep networks: Deep Learning

- Solution: different non-linearities
 - Tanh $\sigma(z) = \tanh(z)$
 - ReLu $\sigma(z) = \max(0, z)$
 - SoftMax $\sigma(z)_j = \frac{e_j^z}{\sum_k e_k^z}$
 - Sum of the output neurons is normed (add up to 1)
 - Useful for classification (for N classes, N output neurons)
 - Output can be interpreted as class probabilities
- Derivatives have larger value range
 - Vanishing Gradient still not impossible
 - Also opposite (exploding Gradient) possible

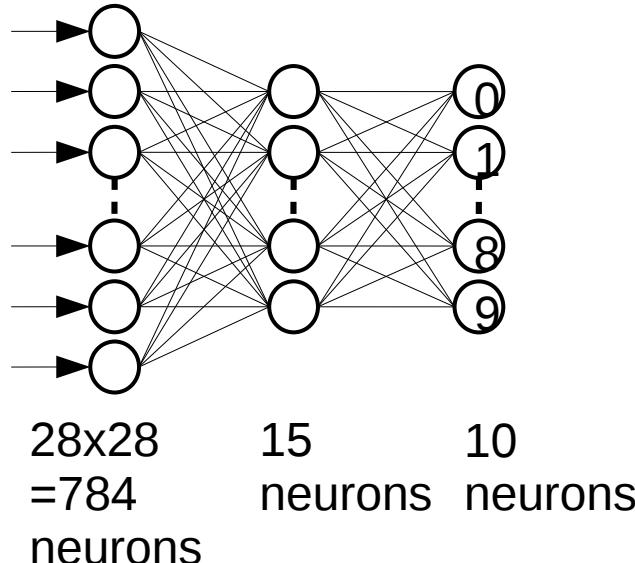
Classification with neural networks

- Example problem: Recognizing handwritten digits
 - MNIST dataset of handwritten digits
 - Images 28x28p
 - Centered on digit



Classification with neural networks

- Recognizing handwritten digits

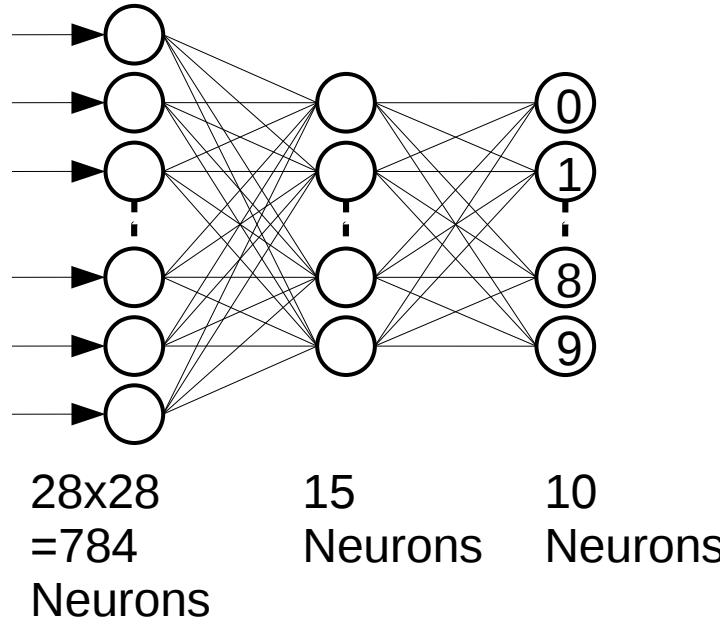


- One output neuron per class: one-hot encoding
- Non-linearity for output neurons: Softmax
 - Since classification problem
- As cost function: Categorical Cross-entropy
 - $C = -\sum_j t_j \cdot \log y_j$

t_j : Groundtruth for class i
 y_j : Prediction for class i

Classification with neural networks

- Recognizing handwritten digits



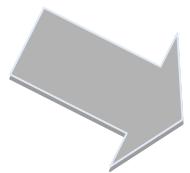
- Each layer extracts feature from the previous layer
 - E.g. Layer 1 could learn to recognize lines, layer 2 could use the lines to recognize shapes (Triangles, ...)
- What happens when a digit that is correctly classified is moved a few pixels?

Example

Input

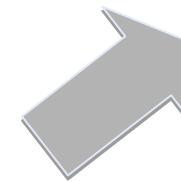
0	2	2	2	0
1	10	2	10	1
2	2	20	2	2
1	10	5	10	1
0	2	10	2	1

Output



$$\frac{1}{10}$$

1	1	1
1	2	1
1	1	1



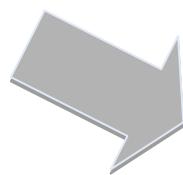
Example

Input

0	2	2	2	0
1	10	2	10	1
2	2	20	2	2
1	10	5	10	1
0	2	10	2	1

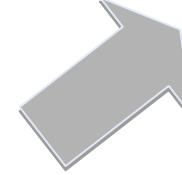
Output

?	?	?	?	?
?	5,1	5,4	5,1	?
?	5,5	9,1	5,5	?
?	6,2	6,8	6,3	?
?	?	?	?	?



$$\frac{1}{10}$$

1	1	1
1	2	1
1	1	1

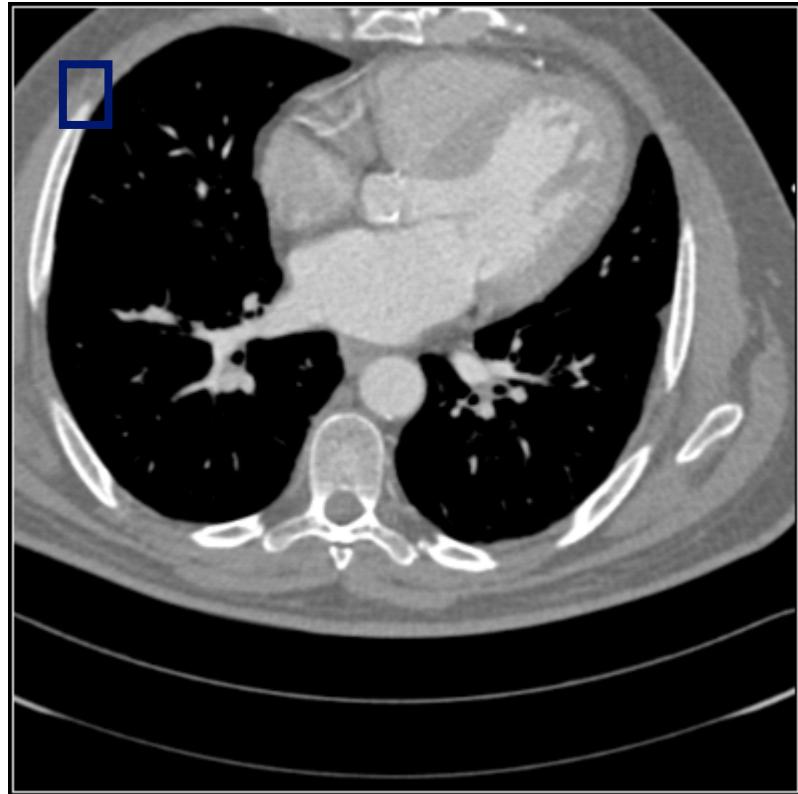
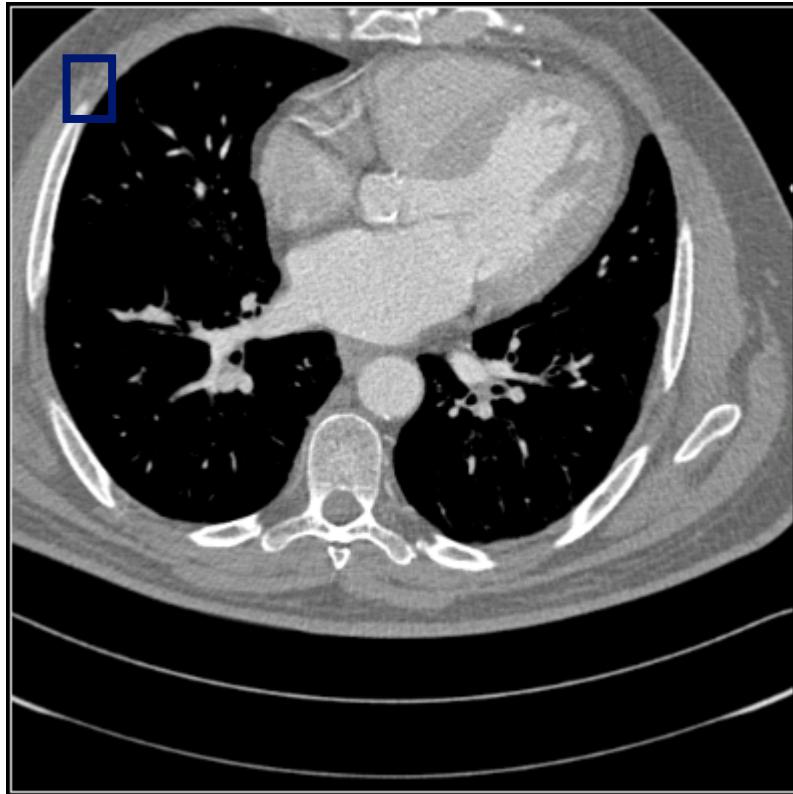


Gaussian filter

- Better smoothing filter than box filter
- Structure is inspired by Gaussian normal distribution
- Normal distribution can be approximated via binomial distribution
- Influence of the environment in dependence to the distance to the image center => no strong flattening of edges

$$\frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}} \cdot \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{y^2}{2\sigma^2}}$$
$$\frac{1}{16} \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \frac{1}{4} \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \cdot \frac{1}{4} \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

Example: Gaussian filter



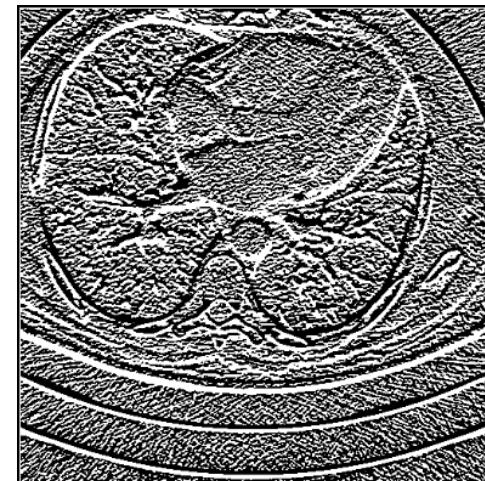
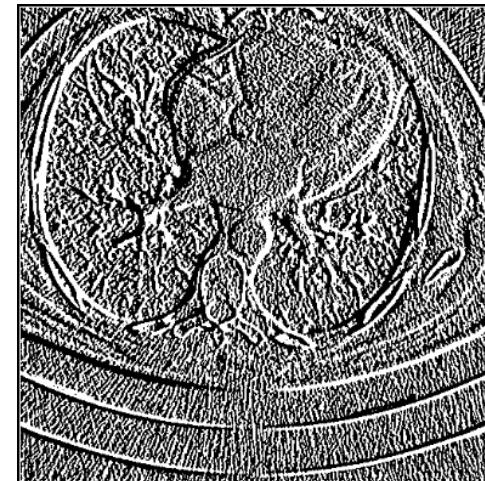
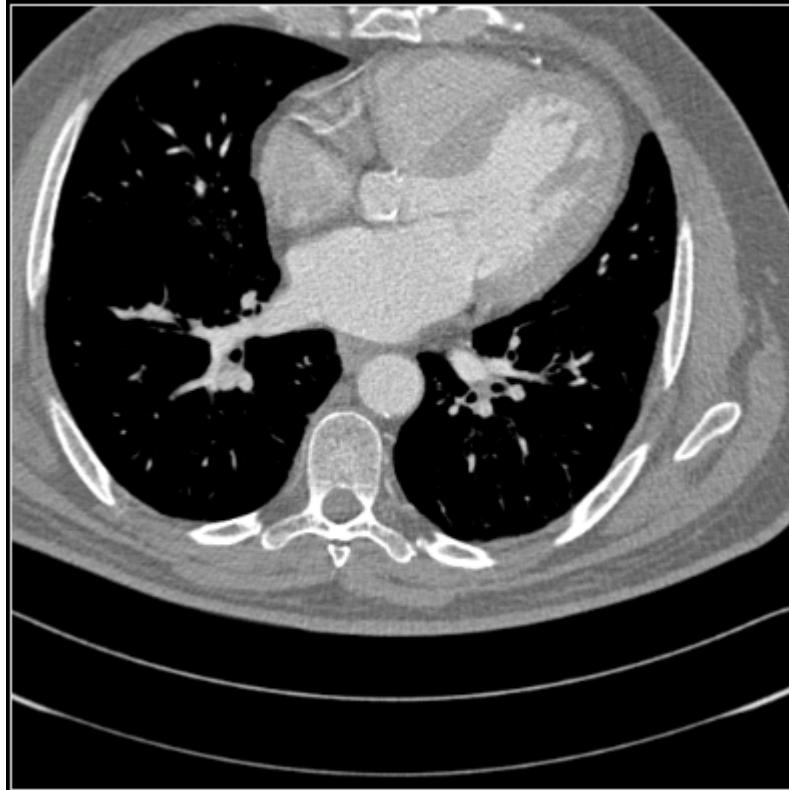
Sobel Filter

- Gaussian-based weighting the difference of the pixel values to enhance edges
- Sobel-X filter enhances vertical, Sobel-Y Filter horizontal edges

$$S_x = \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

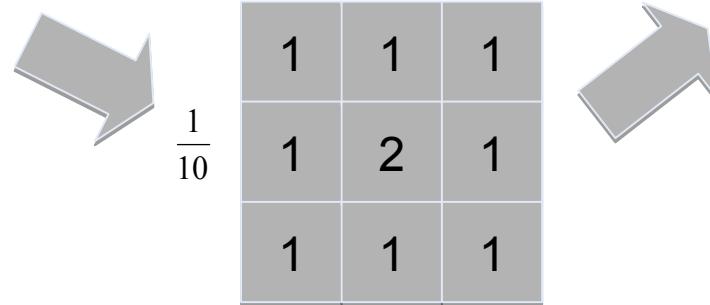
$$S_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Example Sobel Filter



Convolutional Neural Networks (CNNs)

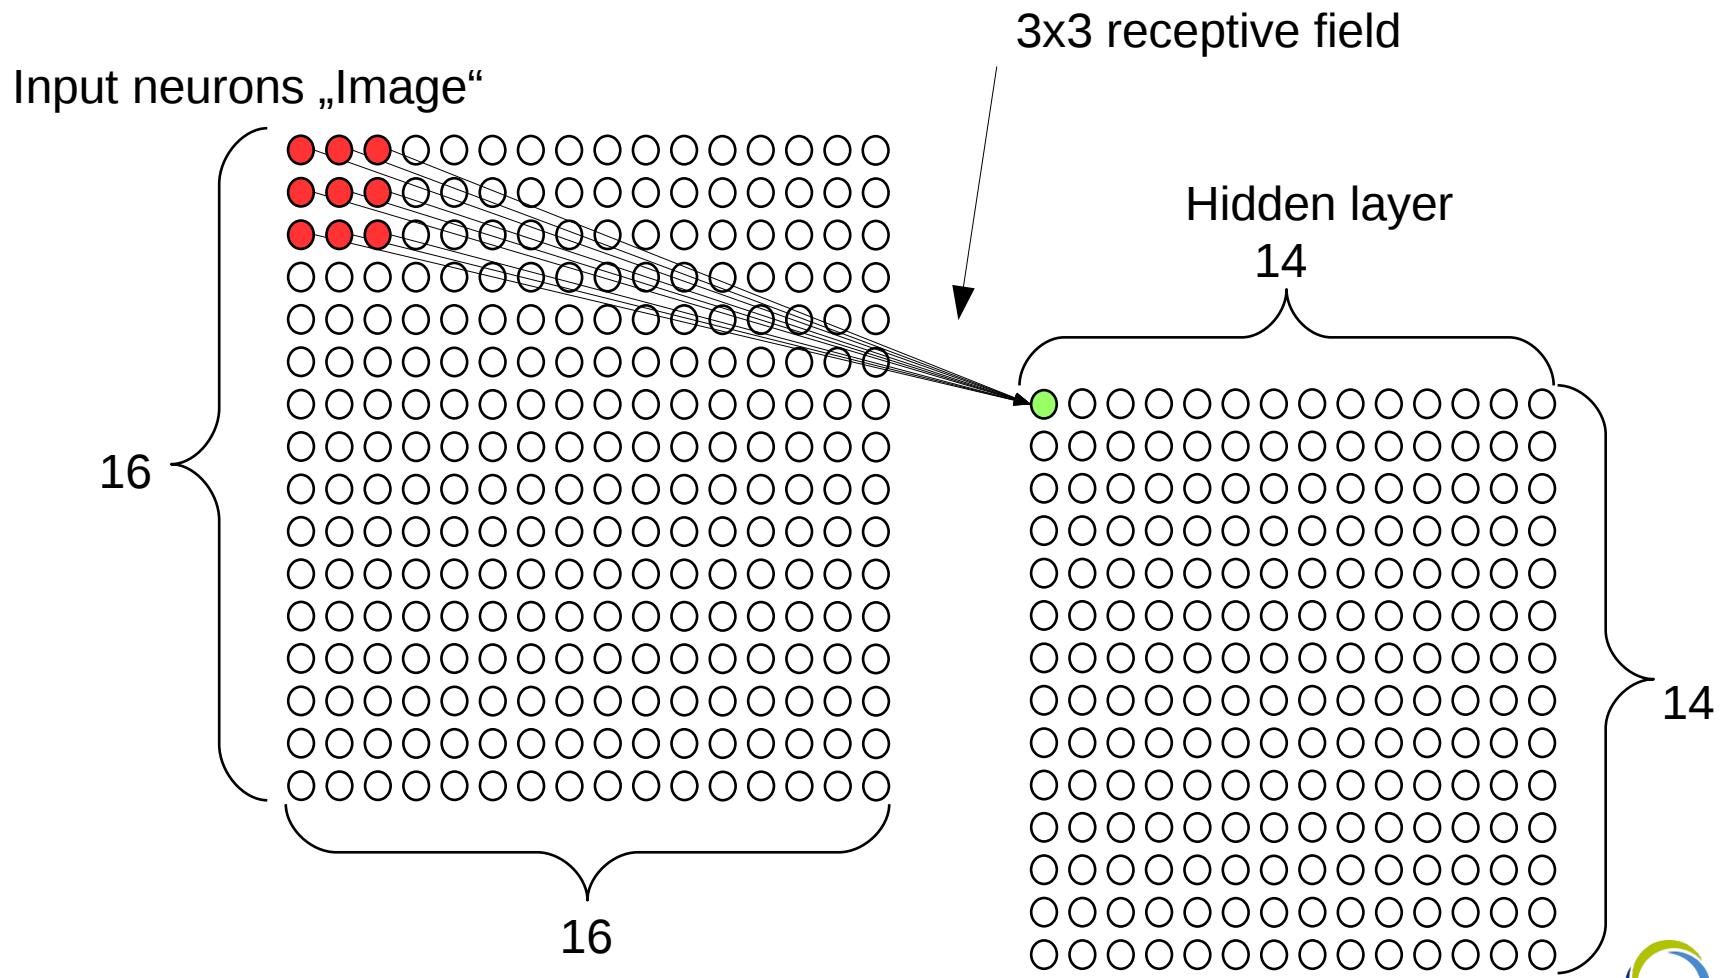
Input					Output				
0	2	2	2	0	?	?	?	?	?
1	10	2	10	1	?	5,1	5,4	5,1	?
2	2	20	2	2	?	5,5	9,1	5,5	?
1	10	5	10	1	?	6,2	6,8	6,3	?
0	2	10	2	1	?	?	?	?	?



$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	$P_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$	$P_y = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$	$S_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$	$S_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$
Gauss	Prewitt			Sobel

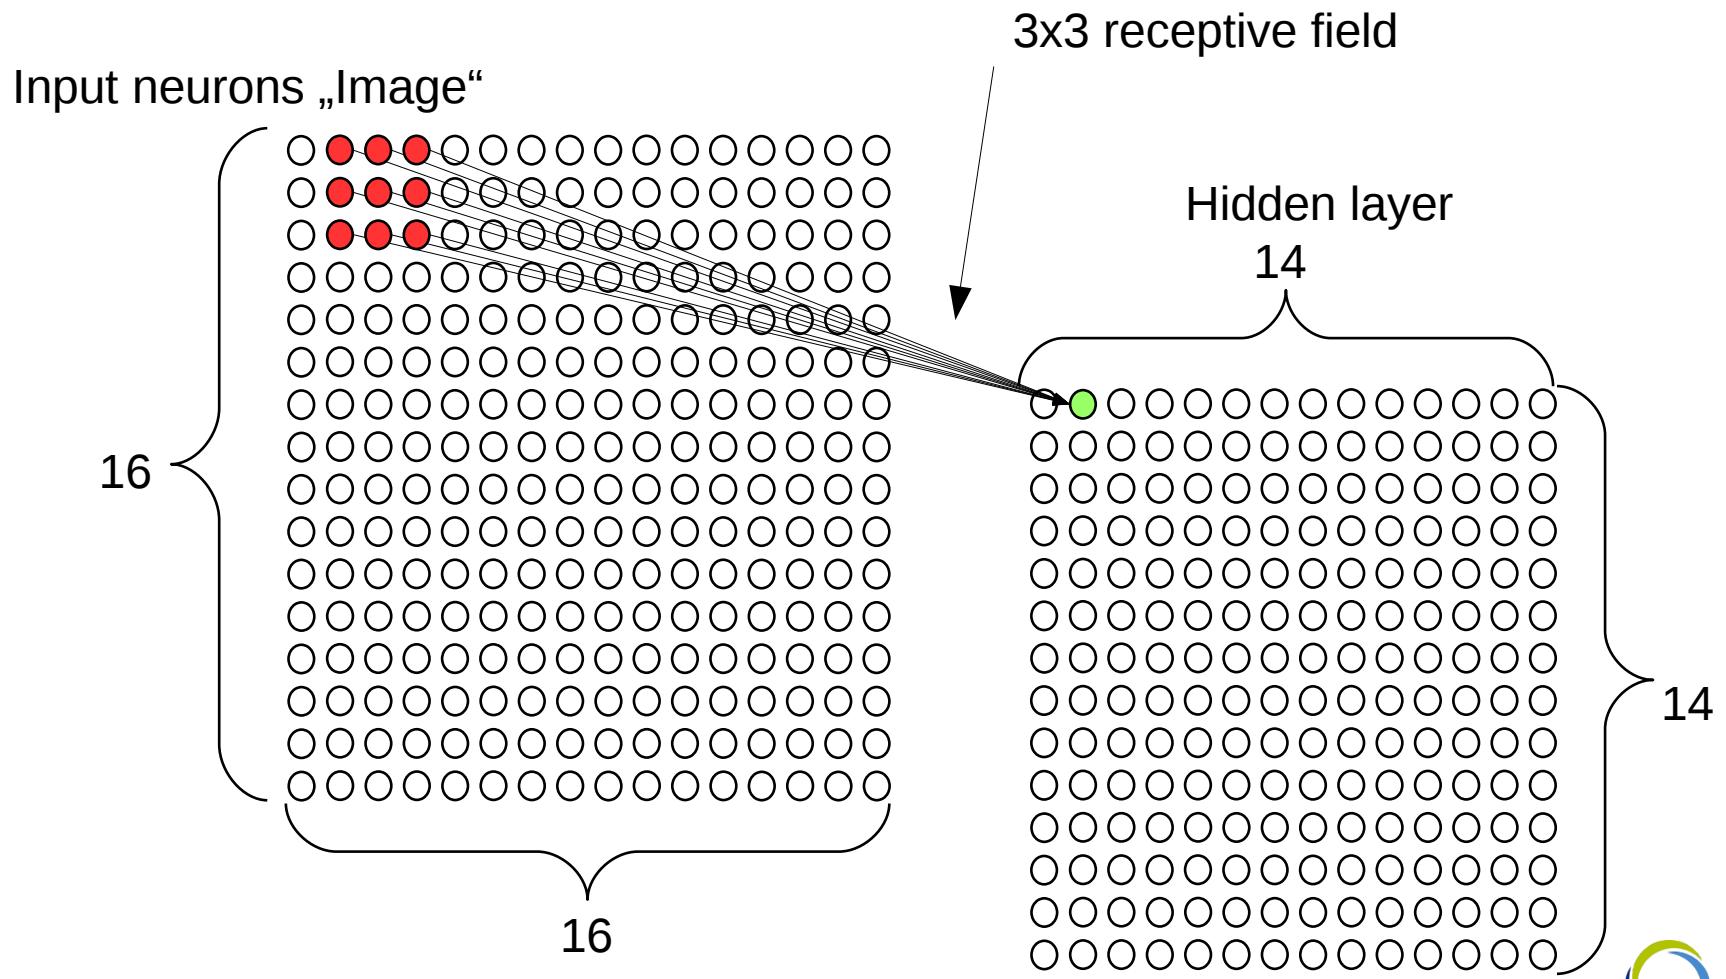
CNNs – Convolutional Layer

- Variant of neural network inspired by filters
 - No complete interconnectedness



CNNs – Convolutional Layer

- Variant of neural network inspired by filters
 - No complete interconnectedness



CNNs – Convolutional Layer

- Variant of neural network inspired by filters
 - No complete interconnectedness
 - Shared weights and biases
 - Combination of weights and bias is called kernel
 - Weight is a 3D-Tensor (# channels x height x width)
 - Output for the j,k^{th} neuron of layer l for a $1 \times 3 \times 3$ kernel

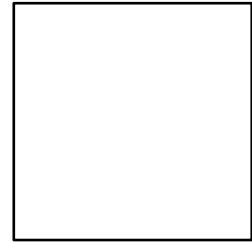
$$a_{l,j,k} = \sigma(b_l + \sum_{m=0}^2 \sum_{n=0}^2 w_{l,m,n} \cdot a_{l-1,j+m,k+n})$$

- All neurons of a receptive field compute the same feature, but on a different position
 - Feature can be found independent of position
- Output matrix is called feature map

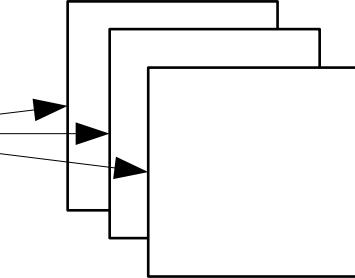
CNNs – Convolutional Layer

- For image processing multiple, varying feature maps are used

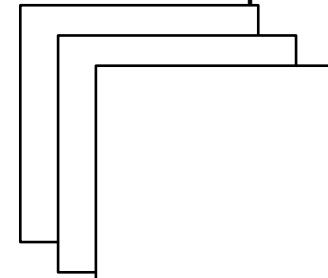
1x16x16 Input



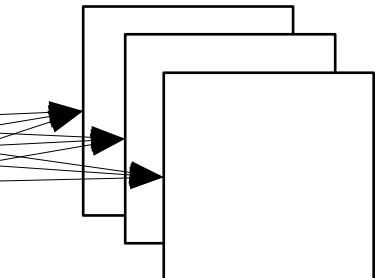
3x14x14 neurons



3x16x16 input



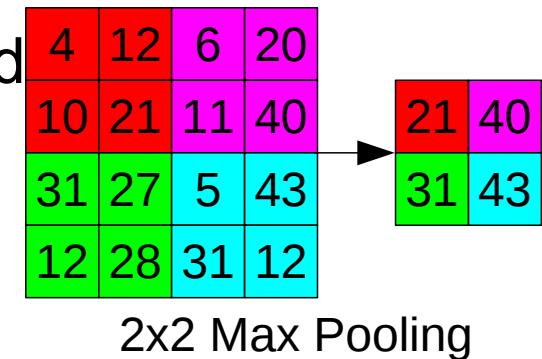
3x14x14 neurons



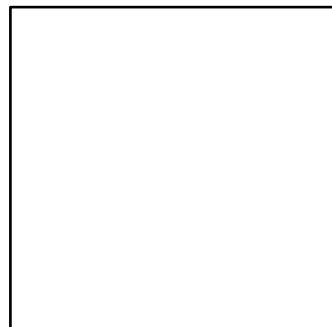
- Feature map for 3x3x3 Kernel has 28 parameters
 - 9 weights per channel: 27, plus 1 bias
 - For 20 feature maps 560 parameters
 - Independent of height and width of input image
 - Compare fully connected: 50372

CNNs – Max Pooling Layer

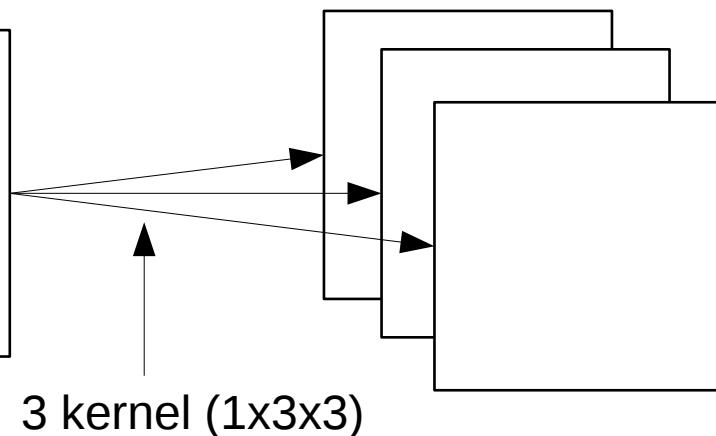
- Reduction of the feature map
 - Retains information if a feature was found
 - Smaller map for following steps
- Selects maximum value from a neighborhood
- Applied to each feature map



1x16x16 input

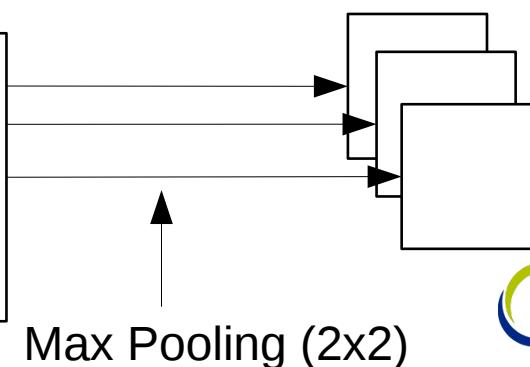


3x14x14 neurons



3 kernel (1x3x3)

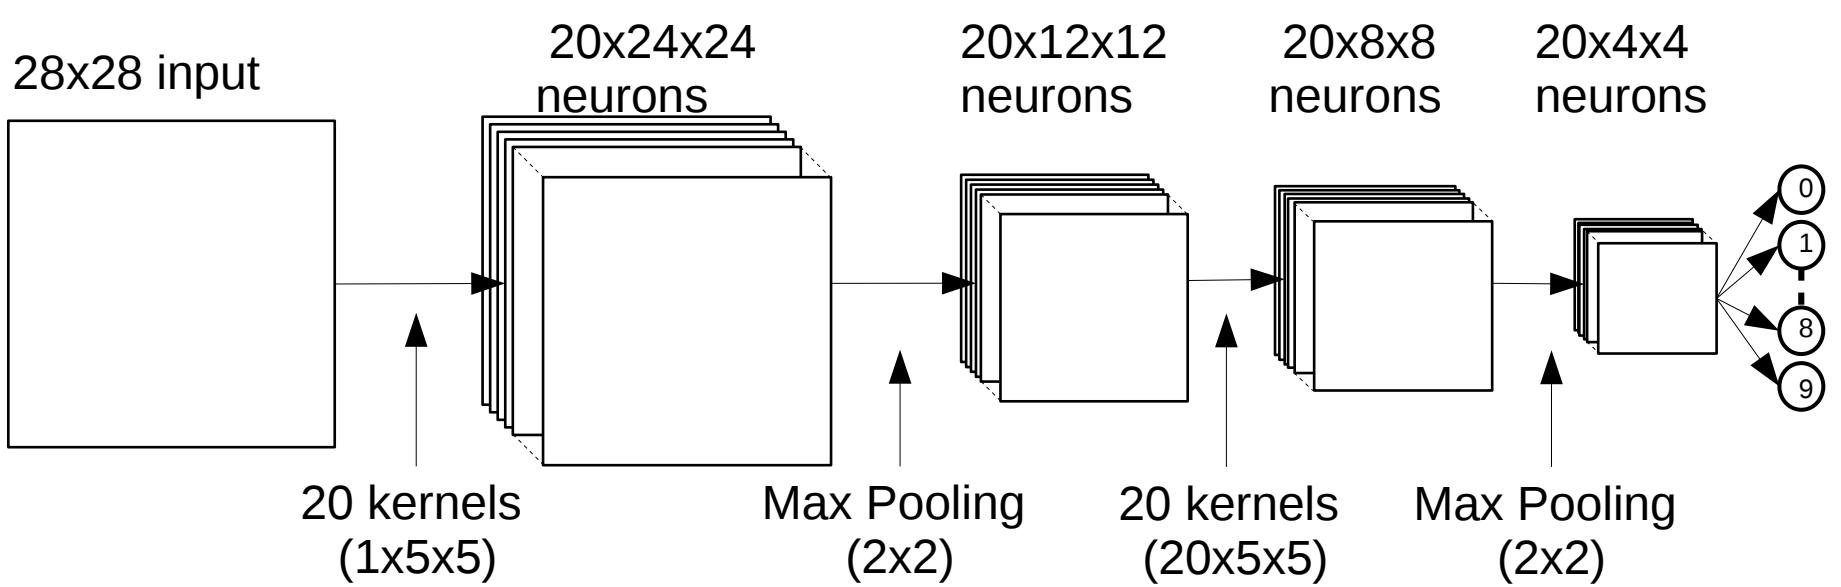
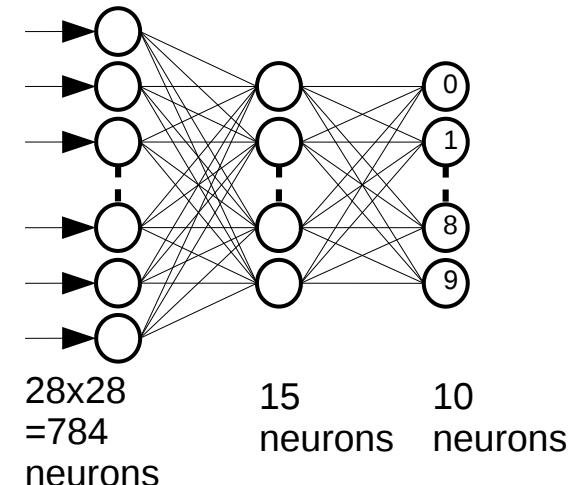
3x7x7 neurons



Max Pooling (2x2)

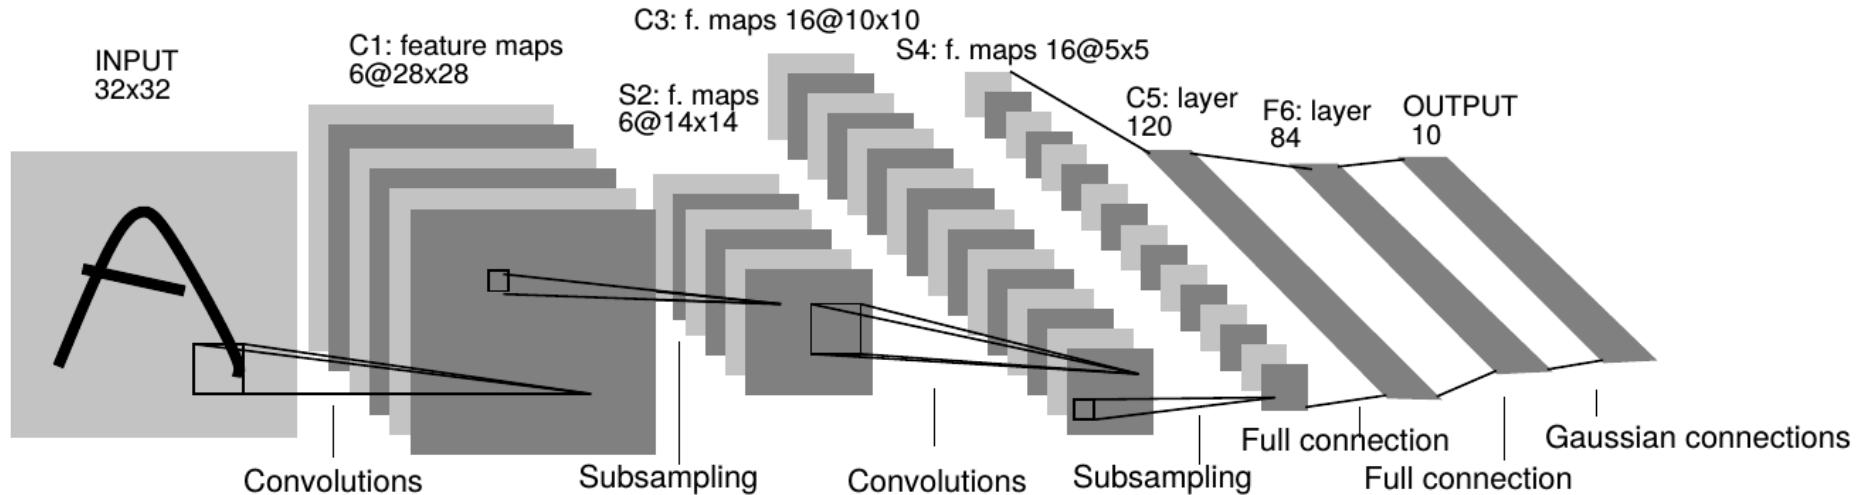
Classification with CNNs

- Recognition of handwritten digits
 - Number parameters
 - Fully connected: 11 935
 - CNN: 13 750



CNN: Example architectures

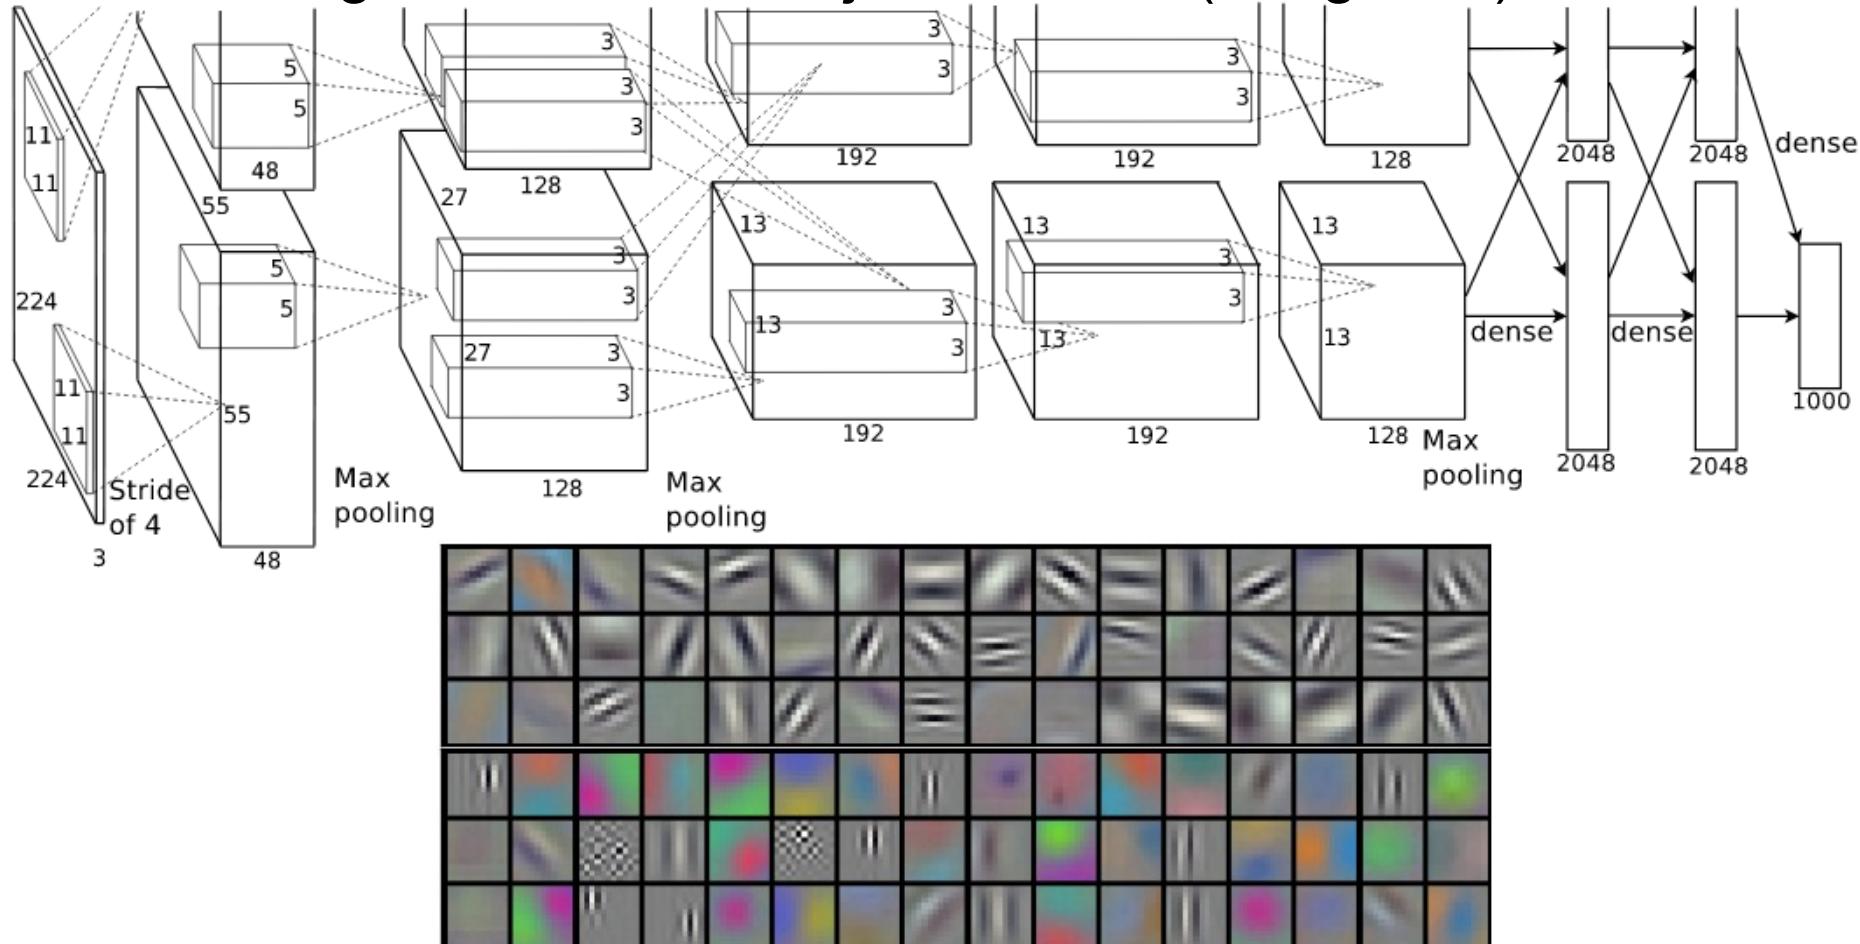
- LeNet: First CNN architectures
 - Recognition of handwritten characters



Y. LeCun, et al., Gradient-based learning applied to document recognition.
Proceedings of the IEEE, november 1998

CNN: Example architectures

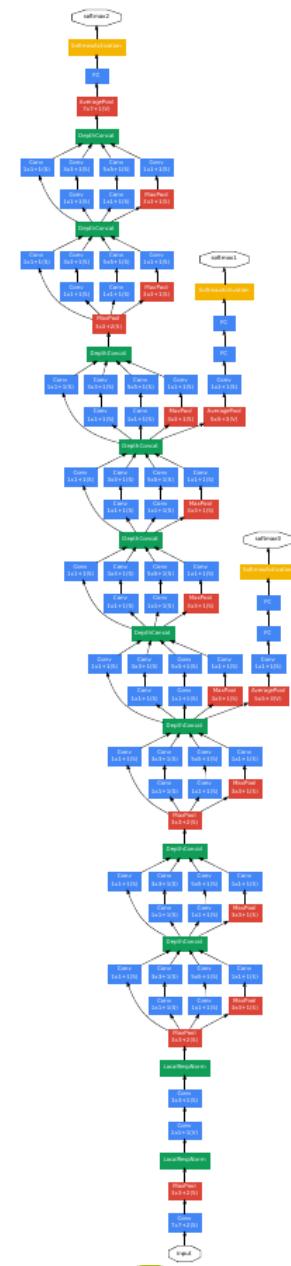
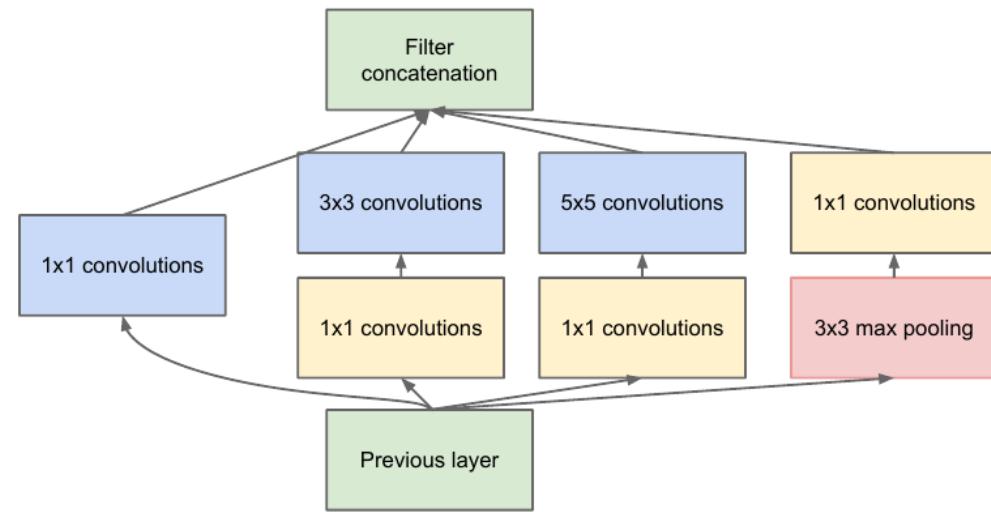
- AlexNet: Popularization of CNNs, first deep CNN
 - Recognition of 1000 object classes (ImageNet)



A. Krizhevsky et al., ImageNet Classification with Deep Convolutional Neural Networks
Advances in Neural Information Processing System, 2012

CNN: Example architectures

- GoogleNet: Especially deep CNN
 - Recognition of 1000 object classes (ImageNet)
 - Only 4M parameters (Compare AlexNet: 60M)

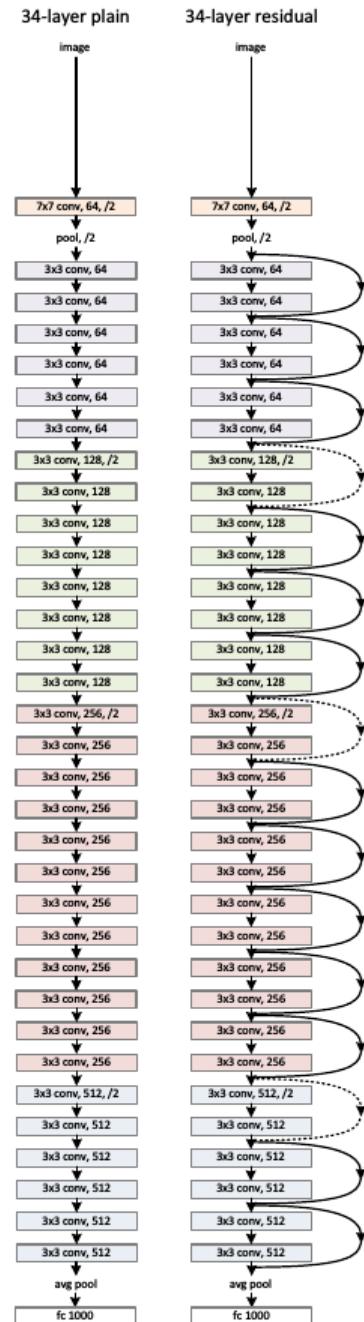
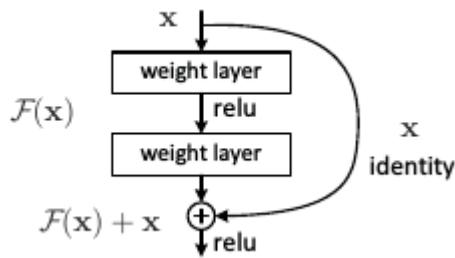


C. Szegedy, et al. "Going deeper with convolutions."

Proceedings of the IEEE conference on computer vision and pattern recognition. 2015.

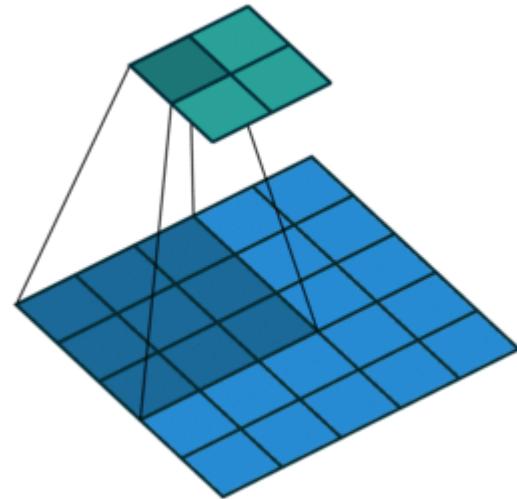
CNN: Example architectures

- ResNet (State of the Art): Even deeper CNNs
 - Recognition of 1000 object classes (ImageNet)
 - Residual Units
 - Skip Chains
 - Allows deep networks with few parameters
 - ResNet32: 0.46M
 - ResNet44: 0.66M
 - ResNet110: 1.7M

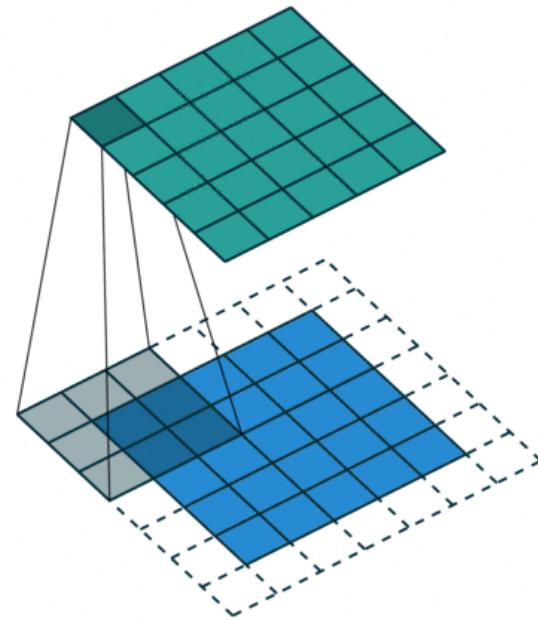


Segmentation with CNNs

- Up to now: Convolution



Convolution with 3x3 Kernel, step size 1, no padding

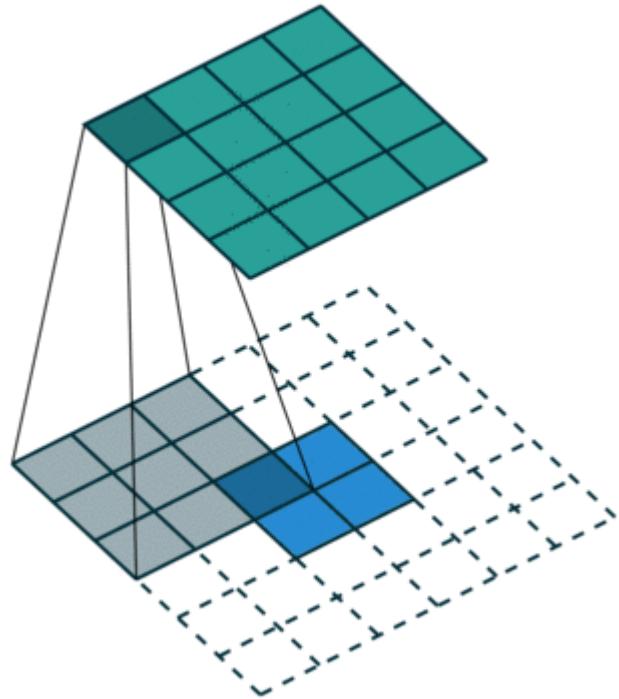


Convolution with 3x3 Kernel, step size 1, padding

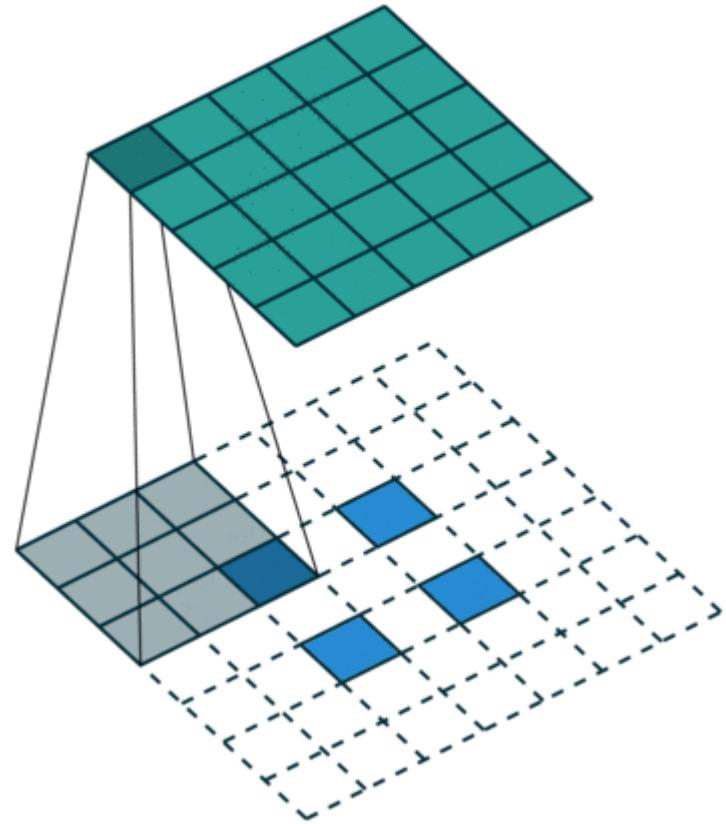
- Resolution is reduced or stays the same
 - Necessary for efficient computation
 - For segmentation output has to have the same size as input

Segmentation with CNNs

- Transposed Convolution



Convolution with 3x3 Kernel, step size 1, padding of 2

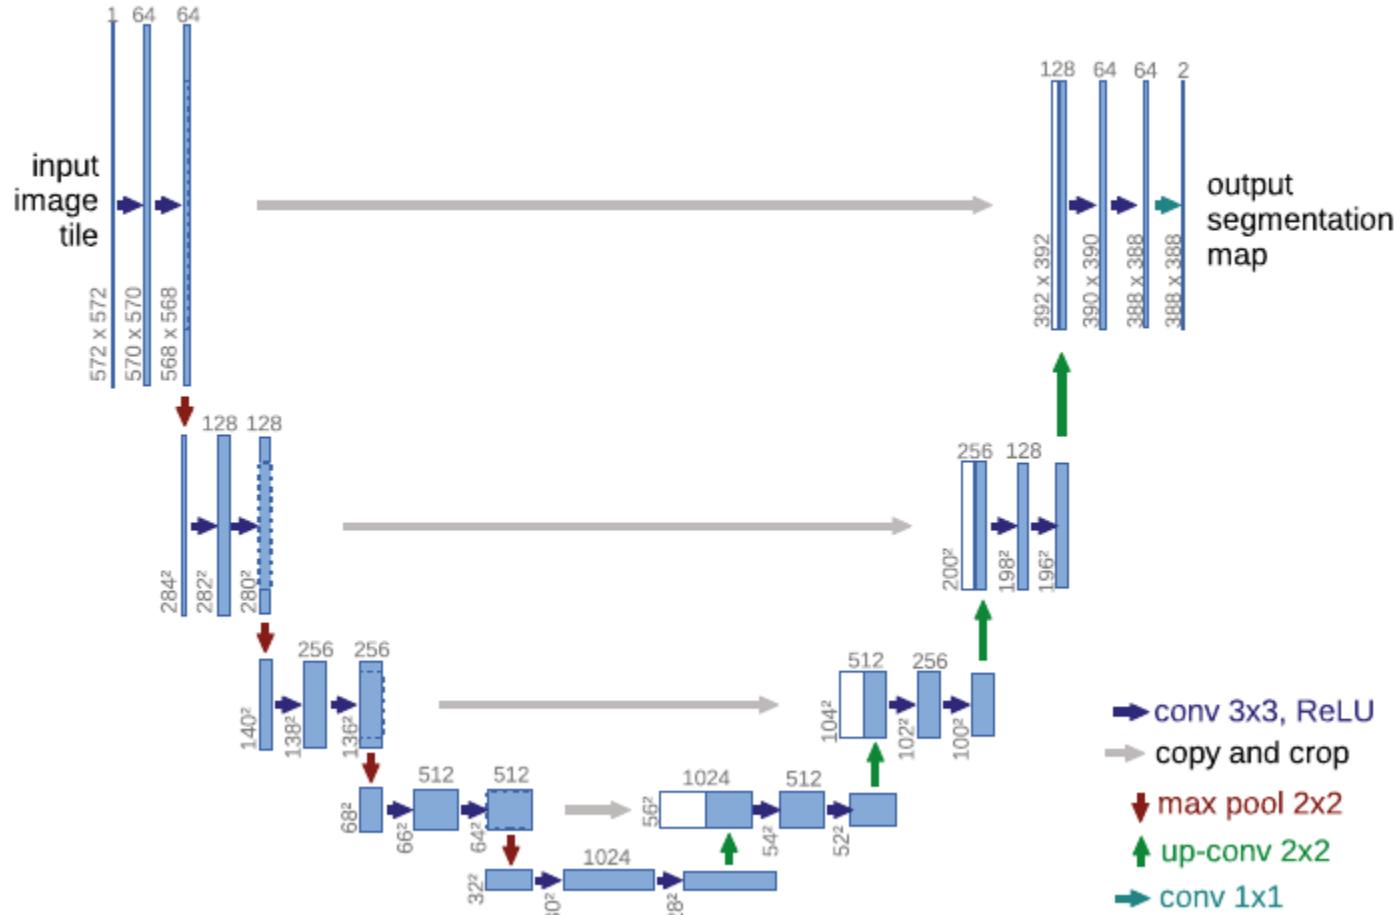


Convolution with 3x3 Kernel, step size 2, Padding

- Can be used to increase resolution

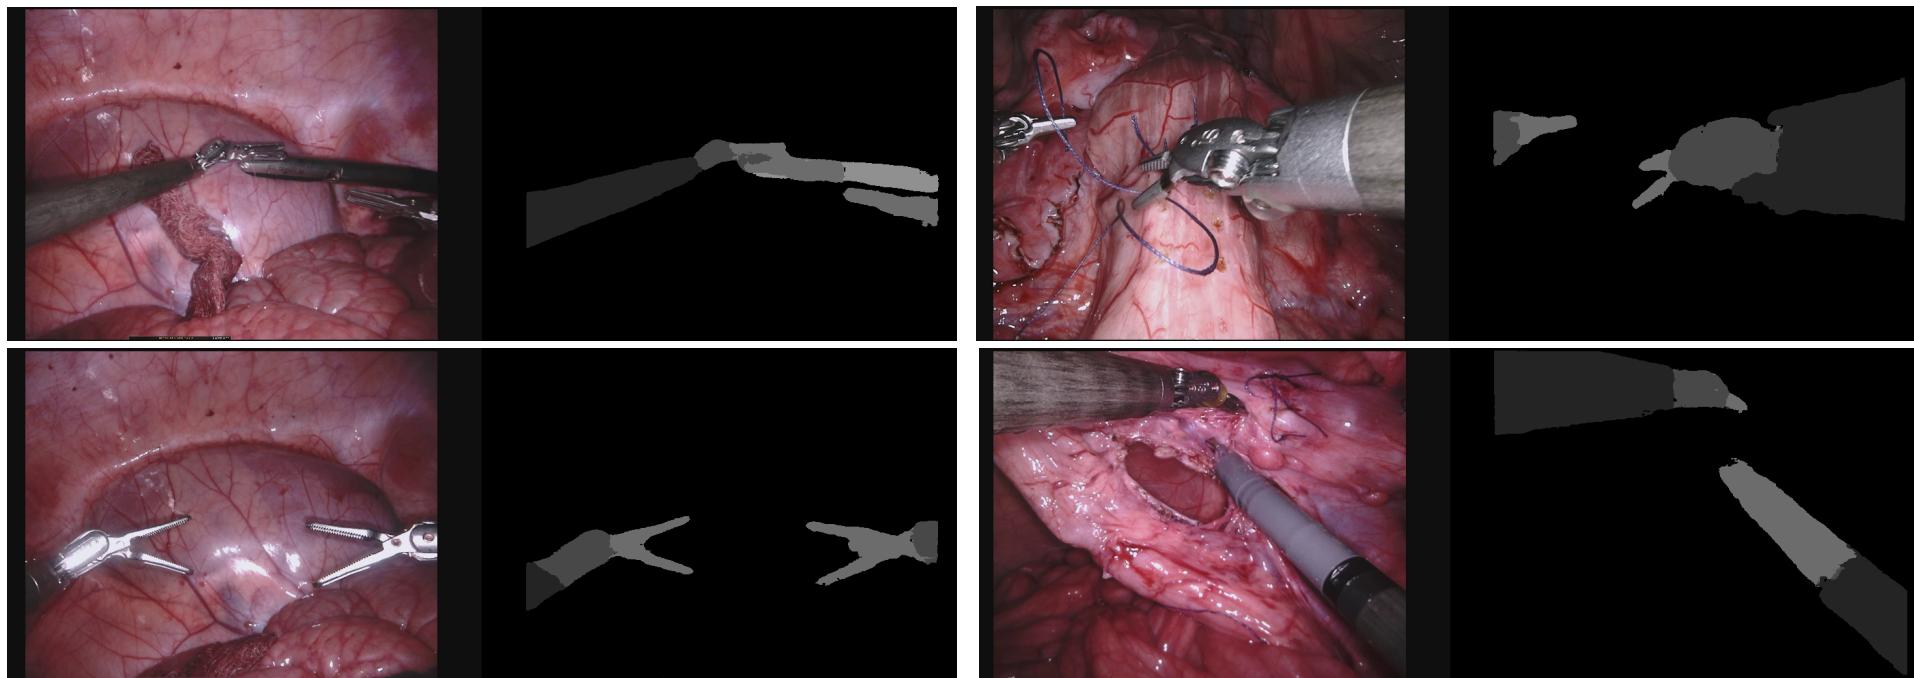
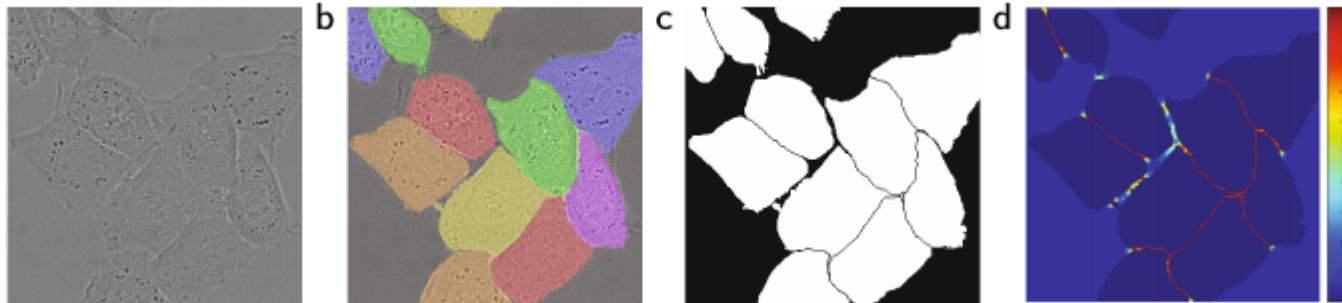
Segmentation with CNNs

- UNet: Pure CNN for segmentation



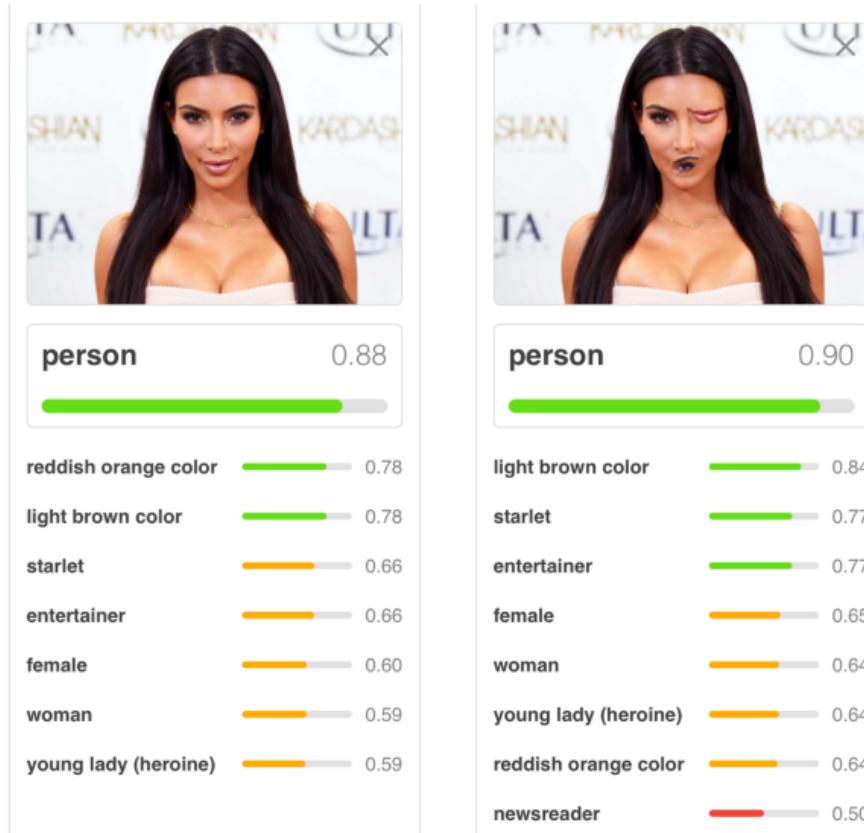
Segmentation with CNNs

- UNet: Pure CNN for segmentation



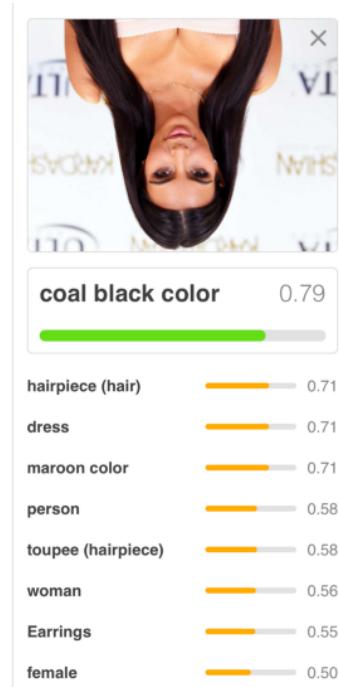
Problem CNNs/Outlook

- CNNs produce State of the Art results for many problems
- Few drawbacks
 - Geometric relations of the features is not explicitly viewed
 - E.g.: A face consists of 2 eyes, a nose and a mouth



Problem CNNs/Outlook

- CNNs produce State of the Art results for many problems
- Few drawbacks
 - Geometric relations of the features is not explicitly viewed
 - Not invariant to rotations



Problem CNNs/Outlook

- CNNs produce State of the Art results for many problems
- Few drawbacks
 - Geometric relations of the features is not explicitly viewed
 - Not invariant to rotations
- Capsule Network
 - Extension of CNNs
 - Saves geometric relations of features
 - Detection via voting
 - E.g.: A face consists of 2 eyes that lie besides each other, a nose below the eyes and a mouth below the nose

Summary

Deep Learning: How can images be segmented and classified using artificial neural networks?

- Perceptron
- Neural Networks
 - Training via Backpropagation
- Convolutional Neural Networks
 - Image classification (what can be seen in an image)
 - Image segmentation (where can something be seen in an image)

Tutorials/Sources

- Great tutorials:
 - <http://neuralnetworksanddeeplearning.com> (very well written)
 - <http://cs231n.github.io/convolutional-networks/>

**Are there
any
questions?**